

Dstl Safe Passage: Detecting and Classifying Vehicles in Aerial Imagery

Vladimir Iglovikov

Physics, PhD
Kaggle Master

Historical overview

December 2016 - March 2017

Kaggle: Dstl Satellite Imagery Feature Detection

Roman Solovyov, Artur Kuzin 2nd place (\$30,000)

Vladimir Iglovikov, Sergey Mushinskiy 3rd place (\$20,000)

- blog posts (rus, eng)
- meetup talks (rus, eng)
- paper (next week)

Organizers spent \$465,000 and got state of the art solutions that they can not use.

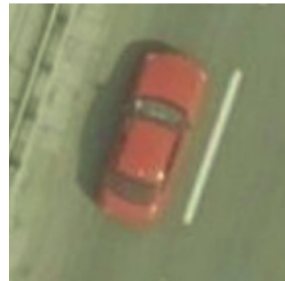
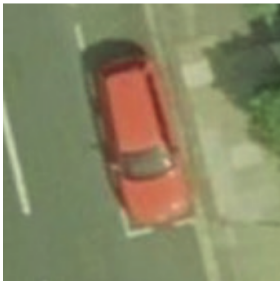
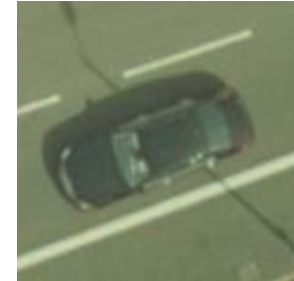
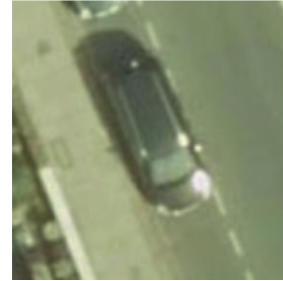
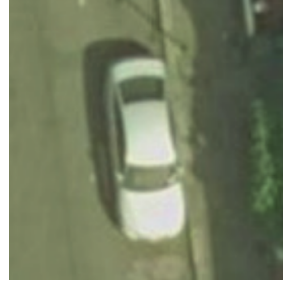
Historical overview

March 2017

- Press release: [Dstl's Kaggle competition has been a great success](#)
- DSTL pays BAE Systems to create their own Kaggle: <https://www.datasciencechallenge.org> and start two competitions (Computer Vision and Natural Language Processing)
- Problems are pretty good, but rules of the competitions are discriminatory (Everyone can participate, but only limited set of people can claim prize money)
- We got verbal and written promise from organizers that rules will be changed.

Problem Statement

- RGB satellite images
- 2000x2000
- 5cm / pixel
- 600 train
- 600 test
- 9 classes



Problem Statement: class distribution

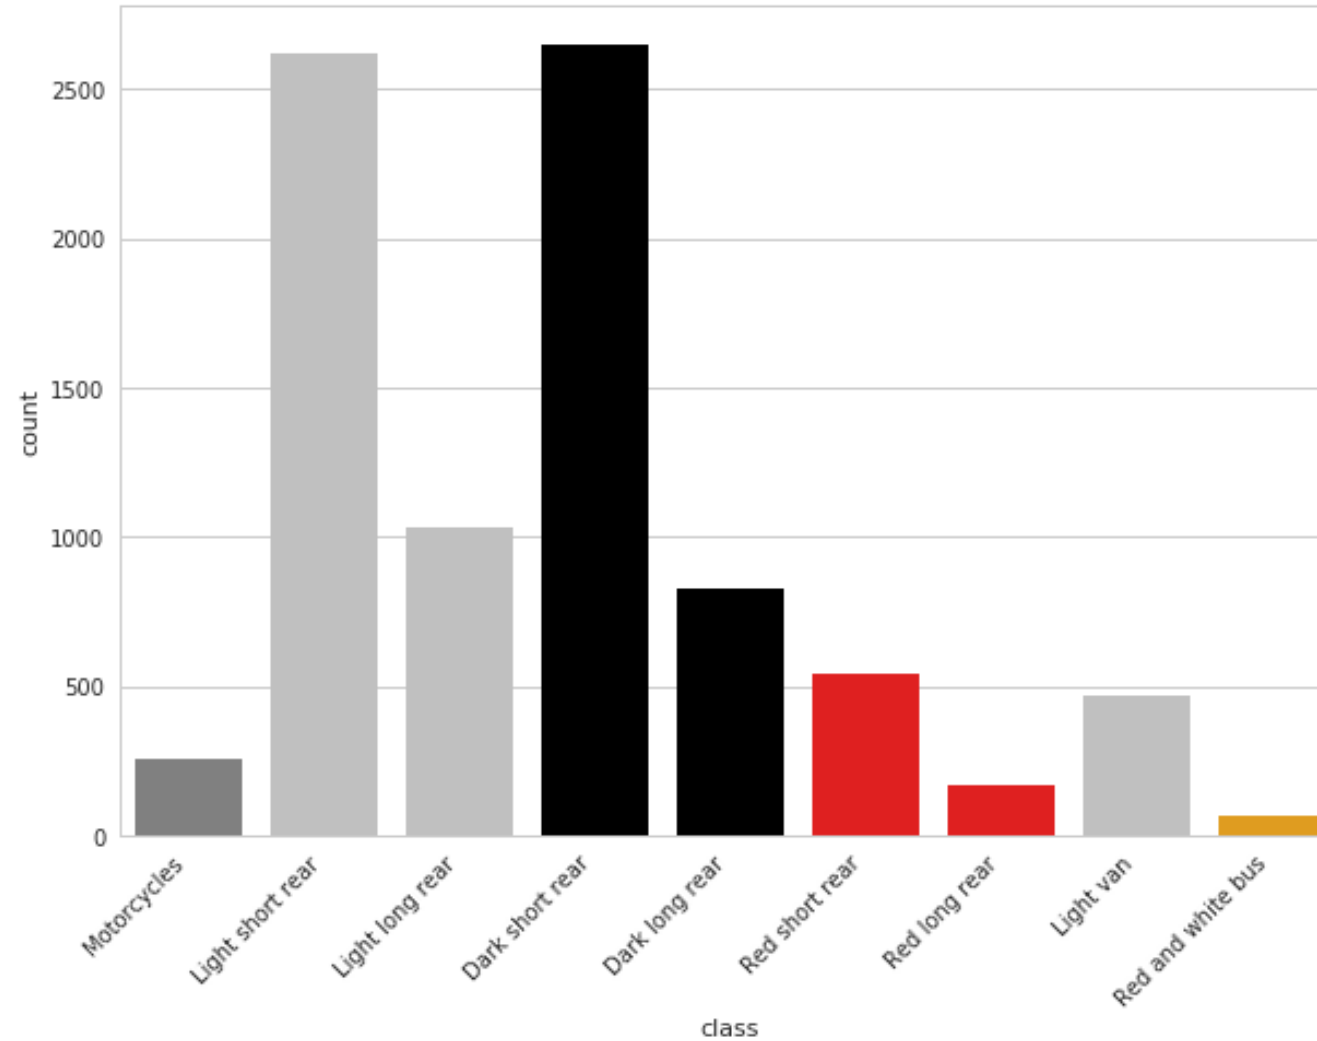


Figure by Vladislav Kassym

Problem Statement

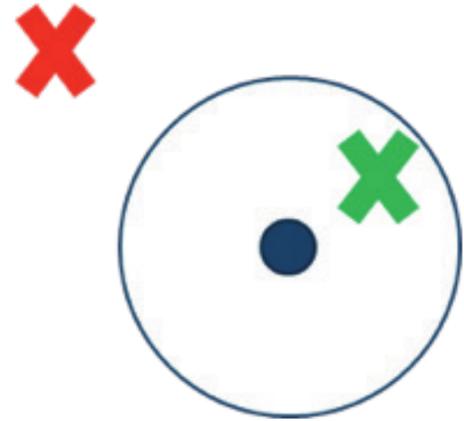


- train: 600 images
- test: 600 images
- 2000x2000
- 5 cm / pixel

One quarter of one image

Evaluation Metric

$$Jaccard = \frac{TP}{TP+FN+FP}$$



Class	Radius
motorcycle	12 pixels (60 cm)
cars	30 pixels (150 cm)
van	40 pixels (200 cm)
bus	45 pixels (225 cm)

Motivation

Why participate?

- Very clean balanced dataset.
- Knowledge in Image Detection.
- Good amount of data. (Not too much, not too little.)
- No data leaks.
- Codebase will be reused in:
 - Kaggle: Cervix
 - Kaggle: Seals
 - ImageNet 2017

Why not participate?

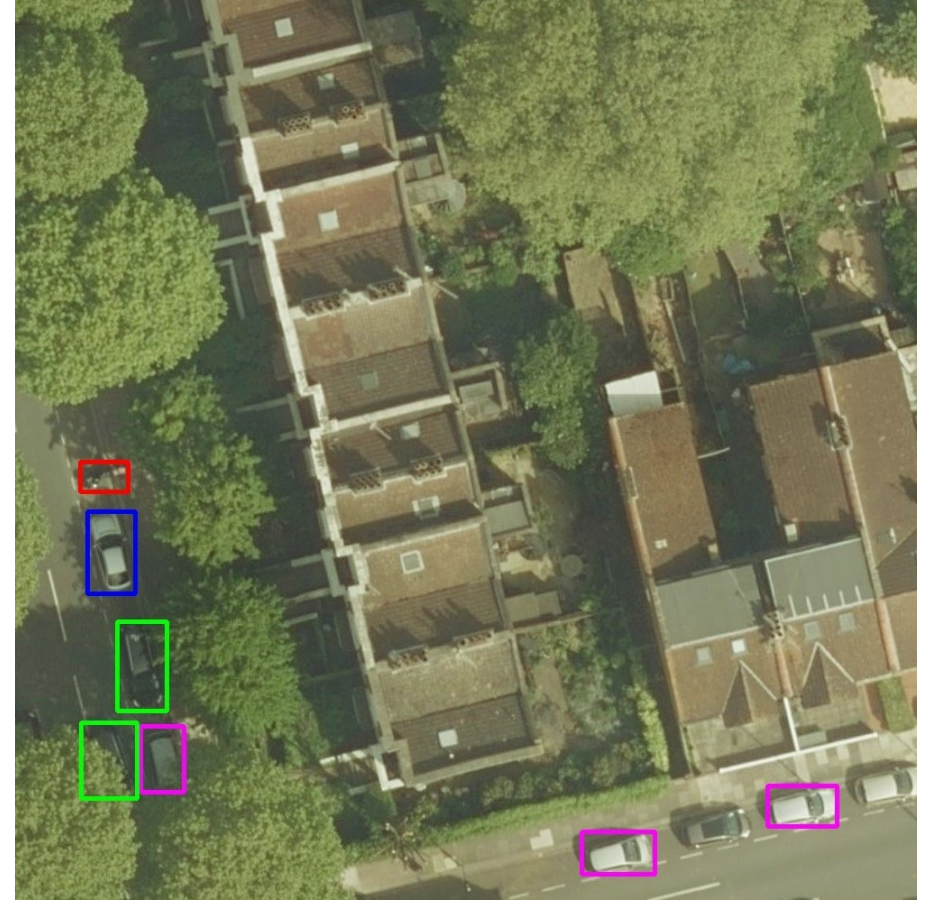
- No way to claim prize money.
- No community.
- Unknown platform. (Hard to sell results.)

Step 1: bounding boxes

Before

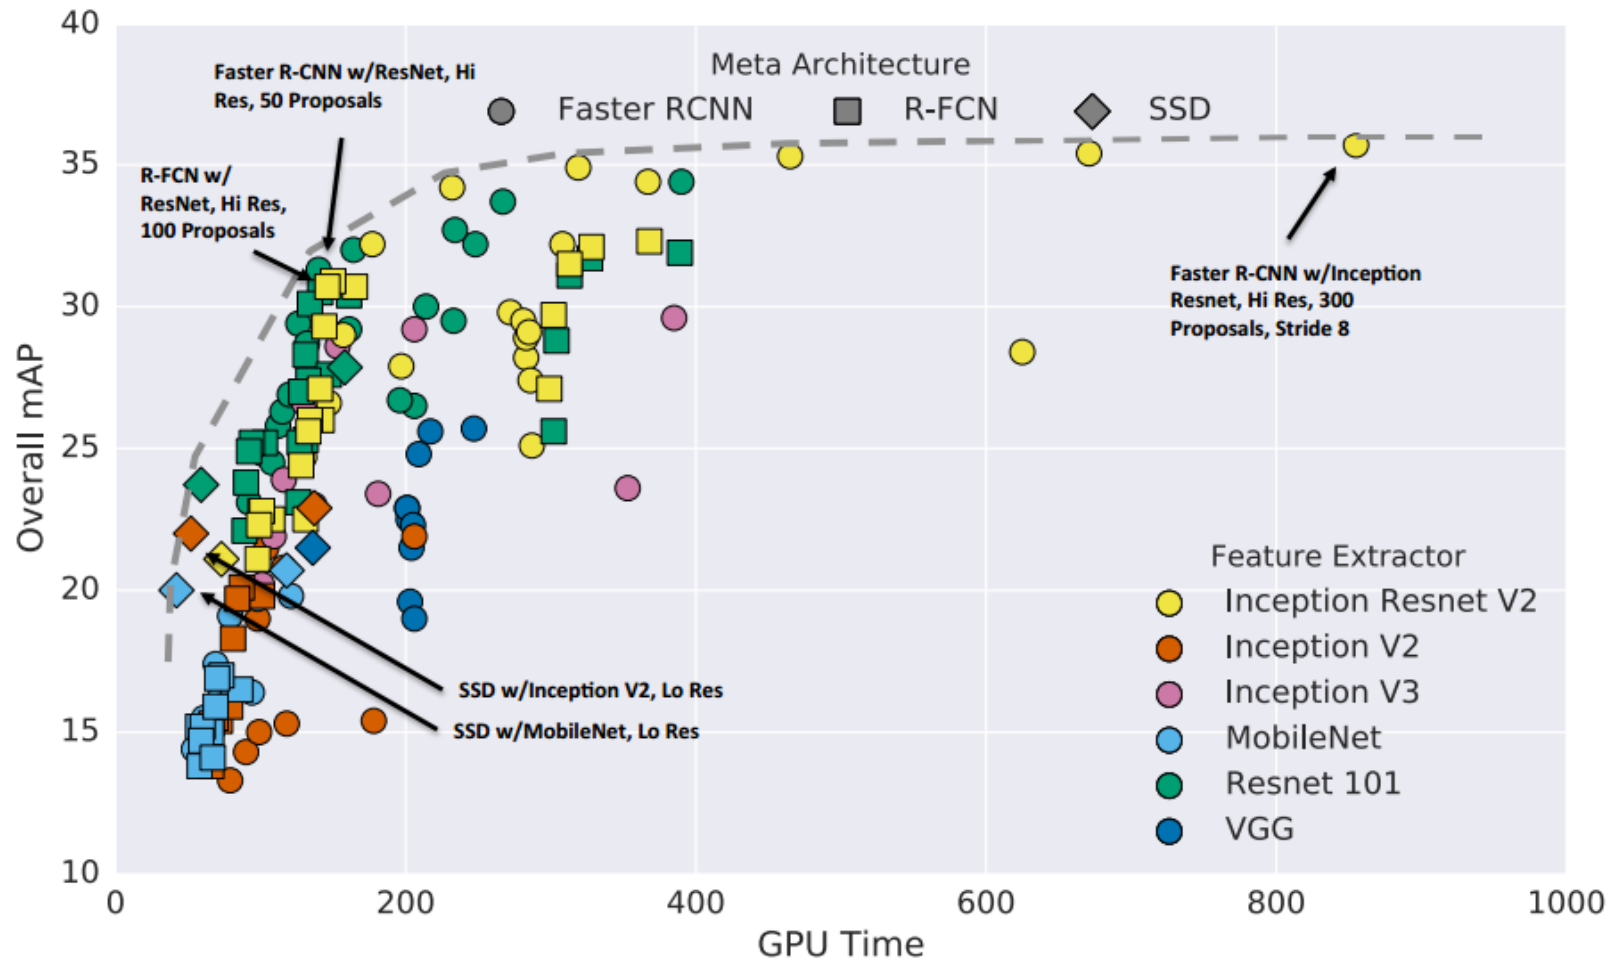


After



~ 10 hours

What network architecture to use?



Speed/accuracy trade-offs for modern convolutional object detectors

What network architecture to use?

Faster RCNN

- Slow to train
- Slow to predict
- Accurate in general
- Accurate on small objects

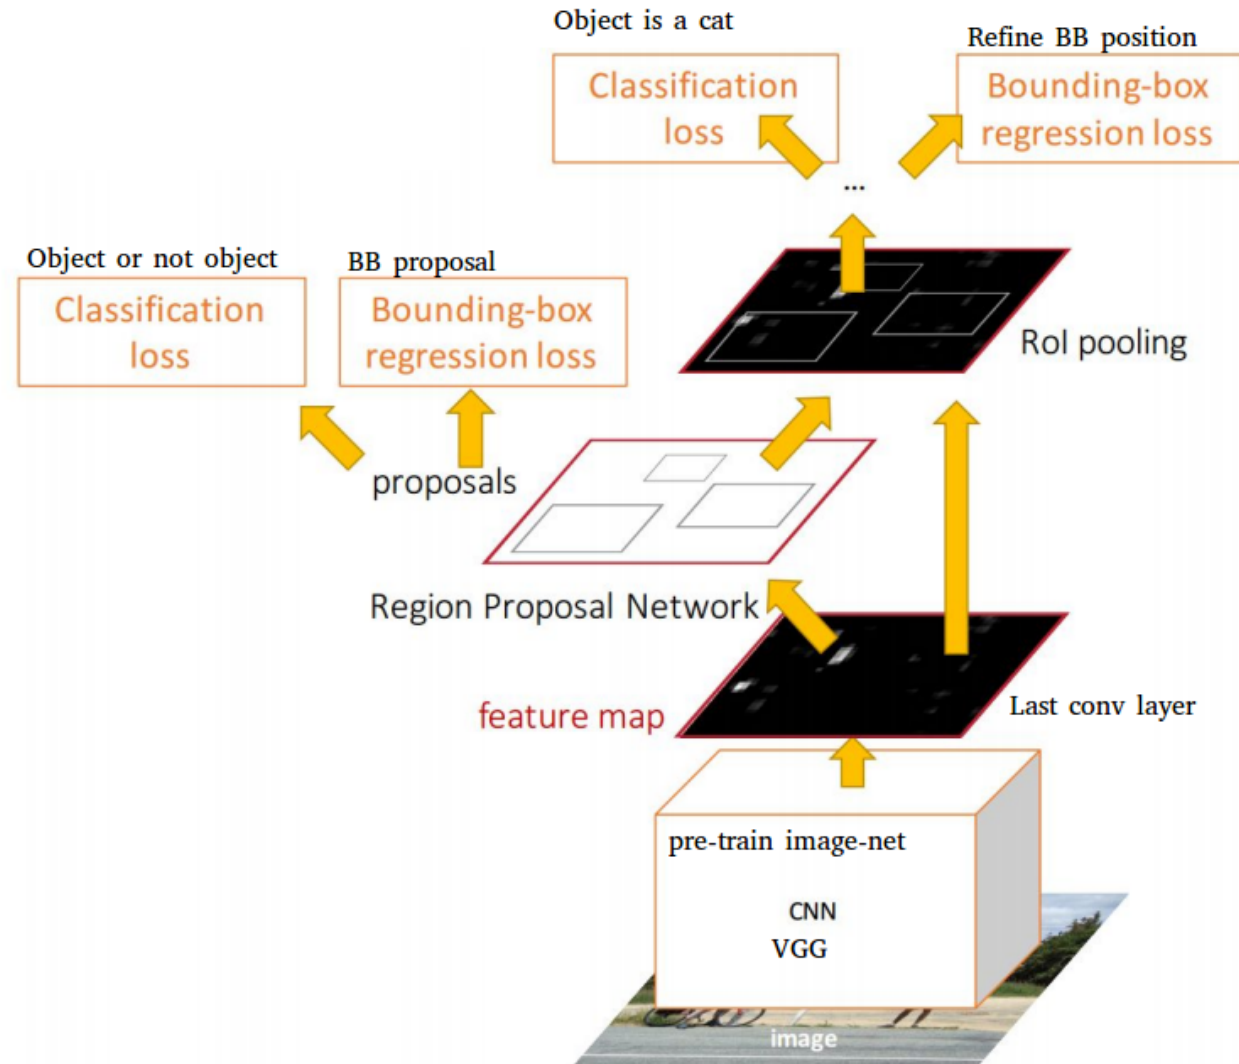
SSD

- Fast to train
- Fast to predict
- Less accurate in general
- Pretty bad with small objects

=>

For this task winner: Faster RCNN

Faster RCNN



What framework to use?

Keras + TensorFlow

- Existing Faster RCNN implementation
- Familiar code base
- Good documentation
- Slow
- Pain to parallelize

MXNet

- Existing Faster RCNN implementation
- Unfamiliar code base
- OK documentation
- Fast
- Zero pain with parallelization

=>

For this task winner: MXNET

Solution

Code - example from MXNet repository

Train

- Faster RCNN + VGG16 base
- random crops 1000x1000
- D4 group augmentation

8 samples/sec

Test

- overlapping tiles
- D4 group augmentation
- Non-Maximum Suppression

20 samples/sec

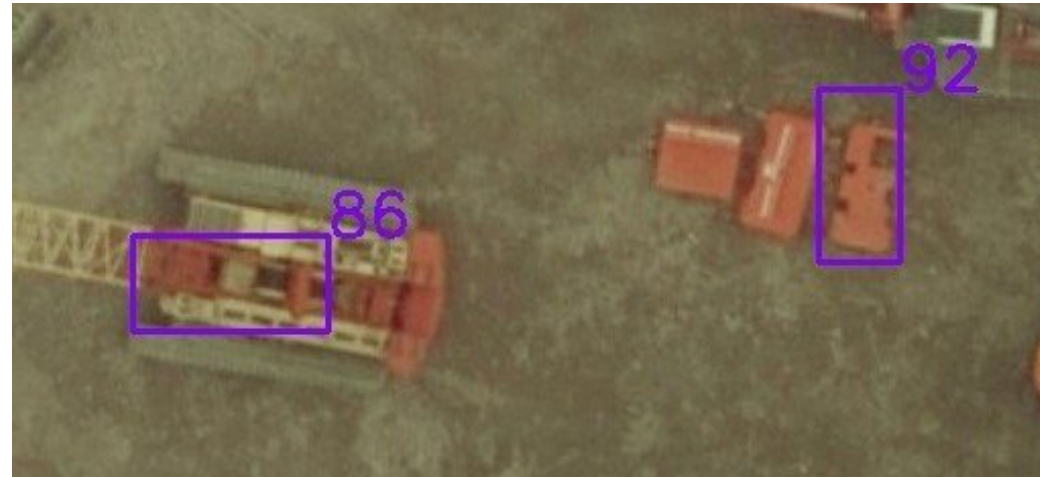
Sources of mistakes: close- packed objects



Sources of mistakes: trains like buses



Sources of mistakes: debris as cars



Main source of mistakes: misclassification

gray car in the shade \Leftrightarrow black car

gray car in the sun \Leftrightarrow white car

blue car in the shade \Leftrightarrow black car

white hatchback \Leftrightarrow white van

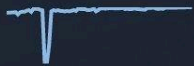





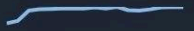



hatchback \Leftrightarrow sedan

\Rightarrow

inconsistent labeling

low predictive power

Top 10 entries

RANK	USER	PUBLIC	PRIVATE	DATE	TREND	ENTRIES
1	gbarbadillo	0.8662	0.8713	17 May 2017, 6:55PM BST		84
2	ternaus	0.8581	0.8569	17 May 2017, 4:47PM BST		31
3	kit1	0.8633	0.8527	16 May 2017, 8:12PM BST		36
4	jane.ostin	0.8490	0.8477	14 May 2017, 8:40PM BST		22
5	Xi_Lian	0.8493	0.8455	16 May 2017, 11:33PM BST		8
6	Kyle	0.8320	0.8369	16 May 2017, 4:05PM BST		48
7	vkassym	0.8333	0.8325	14 May 2017, 8:22PM BST		17
8	dlrocks	0.8142	0.8189	17 May 2017, 10:46PM BST		20
9	cogitae	0.7964	0.8054	17 May 2017, 9:03PM BST		48
10	codewarrior	0.8043	0.8000	04 May 2017, 11:09PM BST		16



Summary

- Centers of cars => bounding boxes (manually)
- Faster RCNN + VGG16, MXnet
- D4 group train and test time augmentation

Hardware

- Intel i7
- 32Gb RAM
- 2 x Titan X (Pascal)

Many thanks to:

- Sergey Mushinskiy
- Vladislav Kassym
- Sergey Belousov