Cdiscount’s Image Classification Challenge

5th place solution

Cdiscount.com - N'

Secure | https://www.cdiscount.com

Cdiscount
N'ÉCONOMISEZ PAS VOTRE PLAISIR.

Rechercher une marque, un produit, ...

Cdiscount services | Cdiscount à volonté | Mon compte Identifiez-vous | Panier 0 article

SOLDES | Electroménager Cuisson Cuisine | Meuble Canapé | Déco Luminaire Linge de maison | Literie Linge de lit | Jardin Brico Animalerie | Informatique | Jeux Video Gaming | TV Son Photo Vidéo | Téléphonie Objet connecté | Auto Moto Mobilité urbaine | Mode Bagage Bijouterie | Bébé Puériculture | Jeux Jouet | Sport Nautisme | Vins Bio Epicerie | Cdiscount Express

Plus que 14h:03m:49:35 pour être livré dès jeudi sur plus de 200 000 articles !

Couette chaude
- 220 x 240 cm
- Anti-acariens

Cdiscount · à volonté

SOLDES -61%

59€99

22€99

SOLDES 2ÈME DÉMARQUE

JUSQU'À 90%**

SUR PLUS DE 3 MILLIONS D'ARTICLES

RETROUVEZ AUSSI NOS OFFRES À VOLONTÉS

TV55" 4K incurvée Smart TV
Classe énergétique A

Haier

Vu dans la presse

699€99

499€99

Produit non soldé

SOLDES 2ÈME DÉMARQUE | MAISON | HIGH-TECH | LITERIE | GROS ÉLECTROMÉNAGER | PETIT ÉLECTROMÉNAGER | JEUX VIDÉO INFORMATIQUE

Commandez votre voiture électrique LYCAR à un prix exceptionnel !
Livrée clé en main, Prête à rouler*

à seulement 9999€ 6999€

3000€ d'économie

*hors assurance

OFFRES À VOLONTÉ
Pc portable gamer 15.6"

OFFRES À VOLONTÉ
Hoverboard + Kit Kart

OFFRES À VOLONTÉ
Canapé d'angle 4 places

OFFRES À VOLONTÉ
TV LED HD 32"

Data

- Almost 9 million products
- More than 12 million images at 180x180 resolution in the train set
- 3 million images in the test set
- 5270 categories
- 1-4 images per product

Some images

1000010653
TELEPHONIE - GPS
ACCESSOIRE TELEPHONE
COQUE TELEPHONE - BUMPER

1000004079
INFORMATIQUE
CONNECTIQUE - ALIMENTATIO
CHARGEUR - ADAPTATEUR SEC

1000004141
INFORMATIQUE
PROTECTION - PERSONNALISA
COQUE - HOUSSE

1000015539
BRICOLAGE - OUTILLAGE - Q
SECURITE MAISON
ALARME AUTONOME

1000010653
TELEPHONIE - GPS
ACCESSOIRE TELEPHONE
COQUE TELEPHONE - BUMPER

1000005744
AUTO - MOTO
PIECES
BOBINE D'ALLUMAGE - BOBIN

1000018290
MUSIQUE
CD
CD MUSIQUE CLASSIQUE

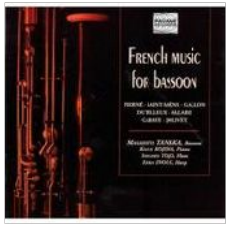

1000010653
TELEPHONIE - GPS
ACCESSOIRE TELEPHONE
COQUE TELEPHONE - BUMPER

1000018306
MUSIQUE
CD
CD VARIETE INTERNATIONALE

1000010961
TV - VIDEO - SON
CASQUE - MICROPHONE - DIC
CASQUE - ECOUTEUR - OREIL

1000015309
LIBRAIRIE
AUTRES LIVRES
AUTRES LIVRES

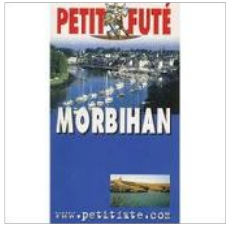

1000010653
TELEPHONIE - GPS
ACCESSOIRE TELEPHONE
COQUE TELEPHONE - BUMPER

Class imbalance

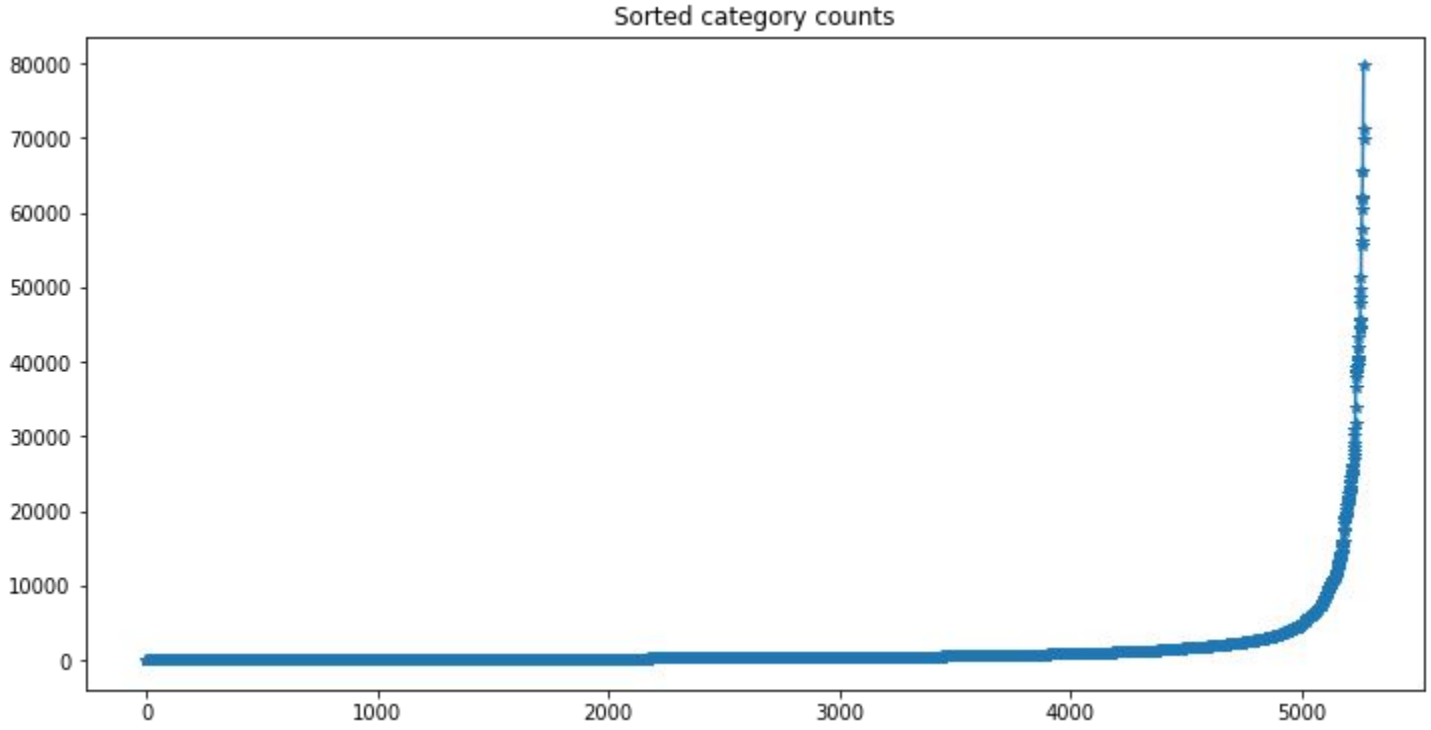

Validation split

- 150k examples for validation
- Strongly correlated with LB score
- Data distribution was the same in both train and test set

Convolutional neural networks

- Resnet models (Resnet50, Resnet101, Resnet151)
- Densenet 161
- InceptionResnetV2

We also tried

- Dual Path Networks
- Xception
- ResNeXt
- SE-Resnet
- VGG
- Nasnet

How to train a deep CNN model faster?

- Use smaller images (160x160 instead of 180x180)
- Don't Decay the Learning Rate, Increase the Batch Size
- Increase augmentation step by step
- Use pretrained weights for initialization

What did not work

- Dropout
- L2 regularization
- Multi-input CNN

How to ensemble several NNs?

- Using weighted mean
- Using weighted geometric mean
- Using 2nd layer model

Ensembling challenges

- Only one model was trained using KFold split (K=4)
- It is impossible to train XGBoost on 5270 classes
- 5270 * 4 features for each product

Hash trick

- One image could be present in the data multiple times
- However, one image could be relevant to several categories
- Finding the best way to use hashes is another challenge

Ensembling approach

- For each product and for each image of this product we extracted TOP-5 probabilities and 5 relevant classes
- We also computed MD5 hash for each picture and matched it with the most common class IDs

After this procedure, we had 80 raw features (40 numerical and 40 categorical)

Ensembling approach

- We added a new feature "possible_class_id" and changed the multi-classification problem to binary classification.
- Next, we were training a model which predicts 'is it true that this product has class_id which is equal to "possible_class_id"'
- We only used samples with possible_class_id which was one of the TOP5 predictions of CNN or one of the possible classes based on hash trick.

L2 models

- XGBoost
- LightGBM
- CatBoost
- Random Forest
- Extra Trees

3rd layer lgbm_2

submit91.csv 0.78577 0.78359

a month ago by Pavel Ostyakov

y_pred_average_before = 0.5 * y_pred_xgboost + 0.1 * y_pred_lgbm_1 + 0.25 * y_pred_lgbm_2 y_pred_sum_after_xgb = 3 * xgb_model1_resnet101_0.75234 + 3 * xgb_model1_incres2_0.75106 + 2 * xgb_model1_lext_resnet_152_0.74686 + 2 * xgb_model1_resnet152_0.74422 + 2 * xgb_model1_resnet50_0.74814 + 1 * xgb_model1_lext_resnet_50_0.72812 + 3 * xgb_model1_zfturbo_4Fold_resnet50_0.75448 + 1 * xgb_model1_densenet161_0.73 y_pred_sum_after_lgbm = 3 * lgbm_model1_resnet101_0.75234 + 3 * lgbm_model1_incres2_0.75106 + 2 * lgbm_model1_lext_resnet_152_0.74686 + 2 * lgbm_model1_resnet152_0.74422 + 2 * lgbm_model1_resnet50_0.74814 + 1 * lgbm_model1_lext_resnet_50_0.72812 + 3 * lgbm_model1_zfturbo_4Fold_resnet50_0.75448 + 1 * lgbm_model1_densenet161_0.73 y_pred_sum_after_lgbm_2 = 3 * lgbm_model2_resnet101_0.75234 + 3 * lgbm_model2_incres2_0.75106 + 2 * lgbm_model2_lext_resnet_152_0.74686 + 2 * lgbm_model2_resnet152_0.74422 + 2 * lgbm_model2_resnet50_0.74814 + 1 * lgbm_model2_lext_resnet_50_0.72812 + 3 * lgbm_model2_zfturbo_4Fold_resnet50_0.75448 + 1 * lgbm_model2_densenet161_0.73 y_pred_sum_after_xgb /= (3 + 3 + 2 + 2 + 2 + 1 + 3 + 1) y_pred_sum_after_lgbm /= (3 + 3 + 2 + 2 + 2 + 1 + 3 + 1) y_pred_sum_after_lgbm_2 /= (3 + 3 + 2 + 2 + 2 + 1 + 3 + 1) y_pred_sum_after = 0.5 * y_pred_sum_after_xgb + 0.25 * y_pred_sum_after_lgbm + 0.25 * y_pred_sum_after_lgbm_2 y_pred_all = y_pred_average_before * 0.6 + y_pred_xgb_multi_new * 0.2 + y_pred_sum_after * 0.3

Product ID and category_id are related

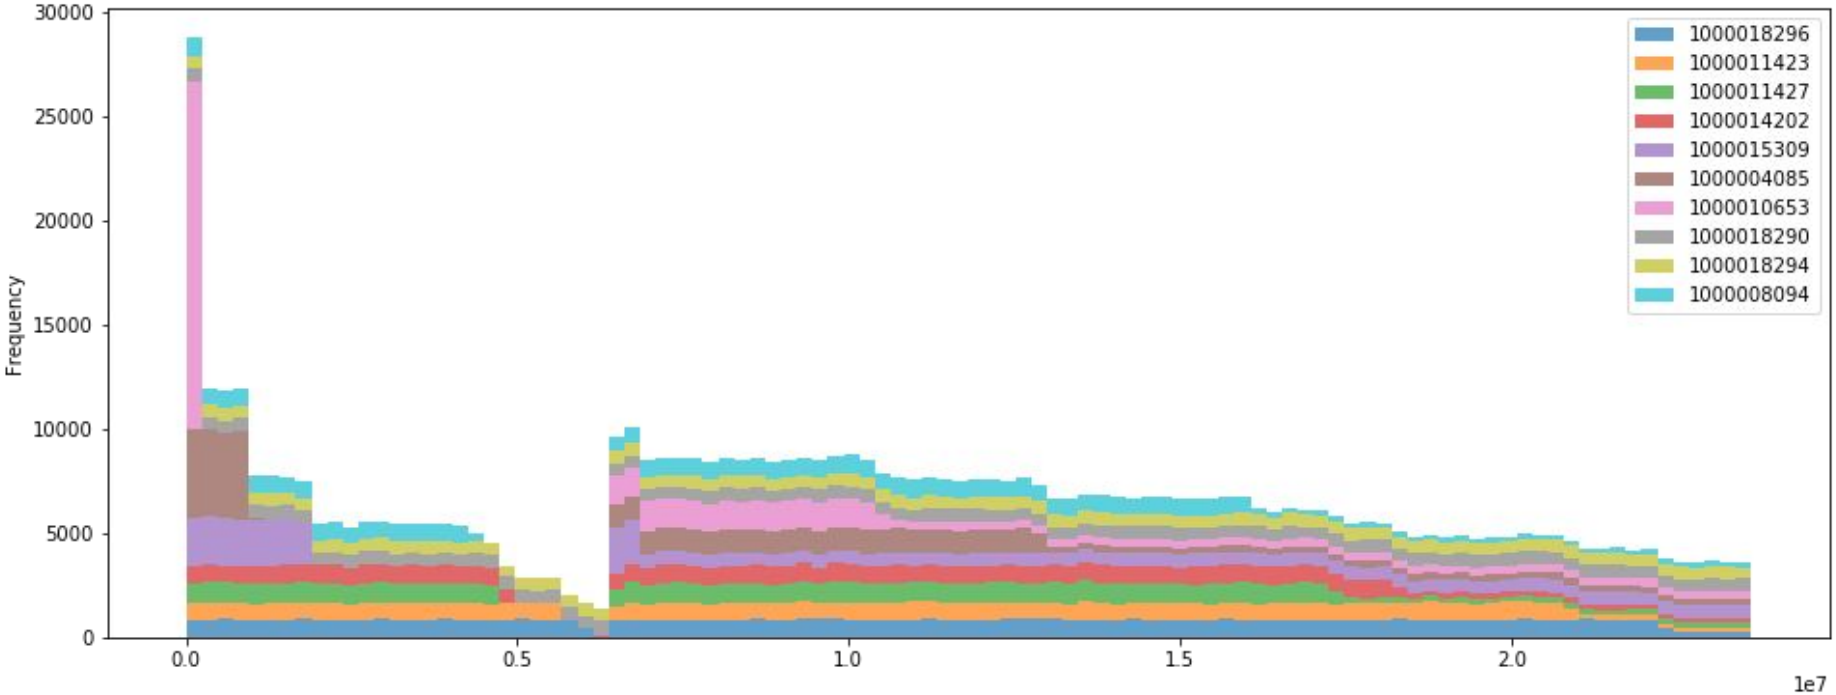

Prior features

For each class_id:

- Frequency of this class in the whole dataset
- Frequency of this class among the closest products based on their IDs (50k closest products).

Third layer features

- Features from CNN
- Predictions of second-layer models
- Priors for categorical features

L3 models

- Three different LightGBM models
- One XGBoost model

It seems that we could improve our score by building a fourth layer model :)

L2 and L3 inference approaches

1. First, average probabilities extracted from CNNs, then build features and use them in second and third layer models.

2. Extract predictions of each CNN, build features, use it for predictions on second and third layers and only then merge predictions of different CNNs.

Semi-supervised learning

- We tried a simple pseudo-label method
- LB score of single Resnet-101 increased from 74.9 to 75.6
- However, it did not improve the score of our ensemble

Our mistakes

- Small validation set
- We did not use out-of-fold validation
- For experiments we trained very large models
- We did not pay enough attention to semi-supervised learning

Software

- ~~Keras~~
- ~~Catboost~~
- Pytorch
- LightGBM
- XGBoost

60	pytorch starter kit here!	Heng CherKeng 3 months ago	last comment by Steven 2mo ago	61
56	Release of trained models (Cdiscount model zoo)	Heng CherKeng 3 months ago	last comment by Nathaniel 5d ago	49
1	how long your model trained?	Five Years 4 months ago	last comment by JuanPizarro 3mo ago	39
5	Single model, single crop Top LB score?	Vinh Nguyen 3 months ago	last comment by Human Analog 3mo ago	35
64	My brief summary of this competition[0.795]	bestfitting a month ago	last comment by bestfitting 5d ago	27
15	batch size trick to accelerate training?	Heng CherKeng 3 months ago	last comment by Vinh Nguyen 2mo ago	25
22	repeating resnet101 LB=0.74 results	Heng CherKeng 3 months ago	last comment by steelrose 2mo ago	22
39	[my ensemble solution]	Heng CherKeng a month ago	last comment by YaGana Sheriff-H... 1mo ago	20
10	Reusing model	Heng CherKeng 3 months ago	last comment by Brian Luo 3mo ago	20
23	recipe for high accuracy model	Heng CherKeng 2 months ago	last comment by tsimins 1mo ago	15
18	Keras starter	Peter Giannakopoulos 3 months ago	last comment by stephenl 3mo ago	14
3	how to use hard samples?	Heng CherKeng 3 months ago	last comment by Heng CherKeng 3mo ago	12
0	automated network design	Heng CherKeng 2 months ago	last comment by Will 1mo ago	8
2	any one trying multiple classifier approach?	Heng CherKeng 3 months ago	last comment by Human Analog 2mo ago	7

Private leaderboard

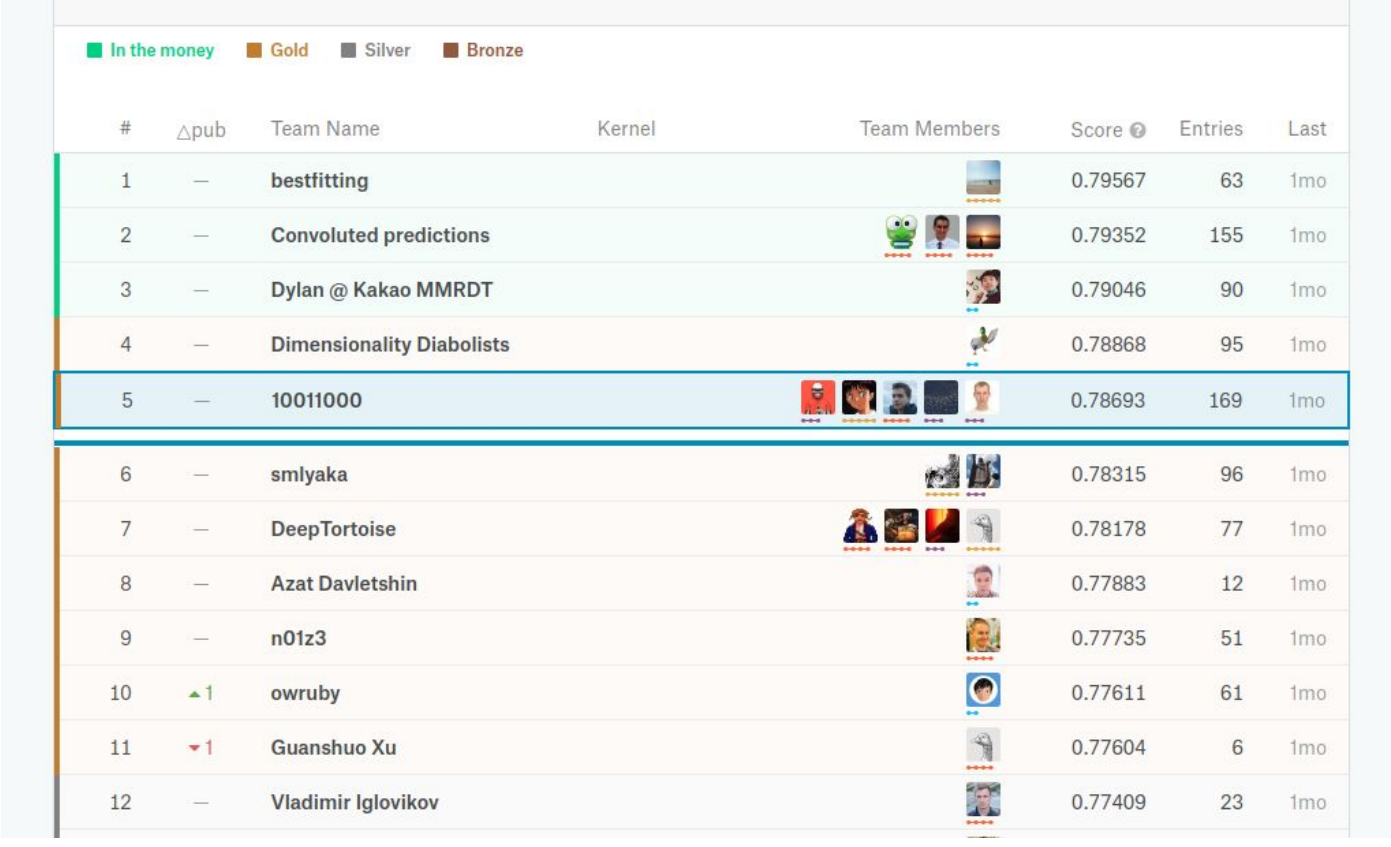

In the money | Gold | Silver | Bronze

#	△pub	Team Name	Kernel	Team Members	Score	Entries	Last
1	—	bestfitting			0.79567	63	1mo
2	—	Convoluted predictions			0.79352	155	1mo
3	—	Dylan @ Kakao MMRDT			0.79046	90	1mo
4	—	Dimensionality Diabolists			0.78868	95	1mo
5	—	10011000			0.78693	169	1mo
6	—	smlyaka			0.78315	96	1mo
7	—	DeepTortoise			0.78178	77	1mo
8	—	Azat Davletshin			0.77883	12	1mo
9	—	n01z3			0.77735	51	1mo
10	▲1	owruby			0.77611	61	1mo
11	▼1	Guanshuo Xu			0.77604	6	1mo
12	—	Vladimir Iglovikov			0.77409	23	1mo

Our team

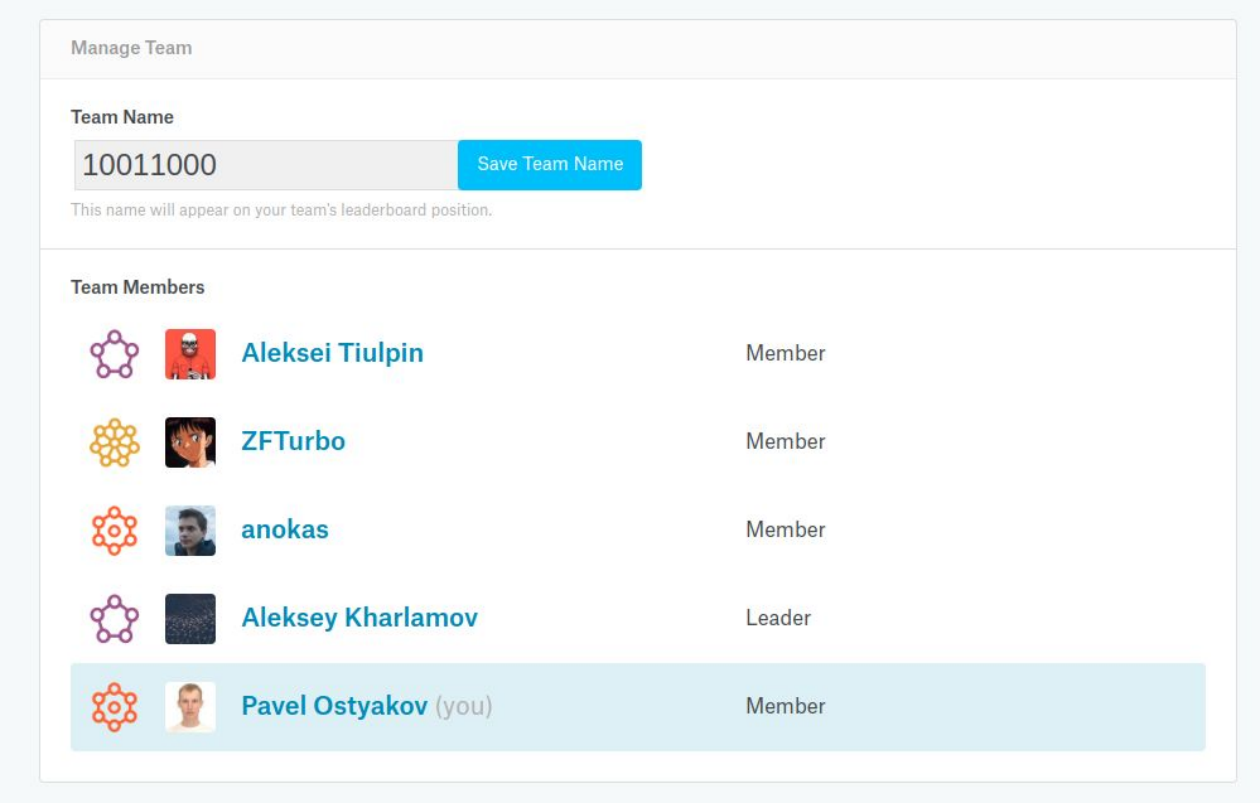
Manage Team
Team Name
10011000
Save Team Name
This name will appear on your team's leaderboard position.
Team Members
Aleksei Tiulpin
Member
ZFTurbo
Member
anokas
Member
Aleksey Kharlamov
Leader
Pavel Ostyakov (you)
Member

bestfitting(1st place)

- Initial model ⟹ ResNet-34
- Representation bottleneck ⟹ 1x1 conv. to increase the channel from 512 to 5270
- x2 FC(5270)
- 11.5 epochs with 160x160 random crop, img. average, get a score >0.72
- ResNet-34 ⟹ ResNet-50
- GPU power + lack of time ⟹ ResNet-101,ResNet-152, Inception-ResNetV2 and InceptionV4

```
lr = 0.0003
if epoch > 7:
    lr = 0.0001
if epoch > 9:
    lr = 0.00005
if epoch > 11:
    lr = 0.00001
```

bestfitting part 2

- Full use of multi-images ⟹ concat. the images into a image
- Split the trainset into 4 parts ⟹ separate models
- Ensemble ⟹ 0.77 LB
- OCR(~0.14% to final sub.) ⟹ ResNet-50 Multi-Input
- Ensemble DenseNet-161,DenseNet-169,DPN-92 + TTA5 ⟹ 0.79 LB

Convoluted predictions(2nd place)

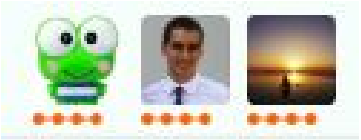

Tricks in training:

1. For base networks used dropout and augmentation=small shift, scale, rotation + flip + rotate90,180,270.Batch_size = 512 to 768, num_epoches = about 20. SDG/learning rate=0.01,0.001,0.0001 (manually tuned). Time to train (4x1080Ti) = 5 days per base networks
2. For Fcnet3 used dropout and only augmentation=flip (due to lack of space in SSD drive). Batch_size = 4096, num_epoches = about 100. SDG/learning rate=0.01,0.001,0.0001 (manually tuned). Time to train (4x1080Ti) = 1.5 days

[Model]

combine.0

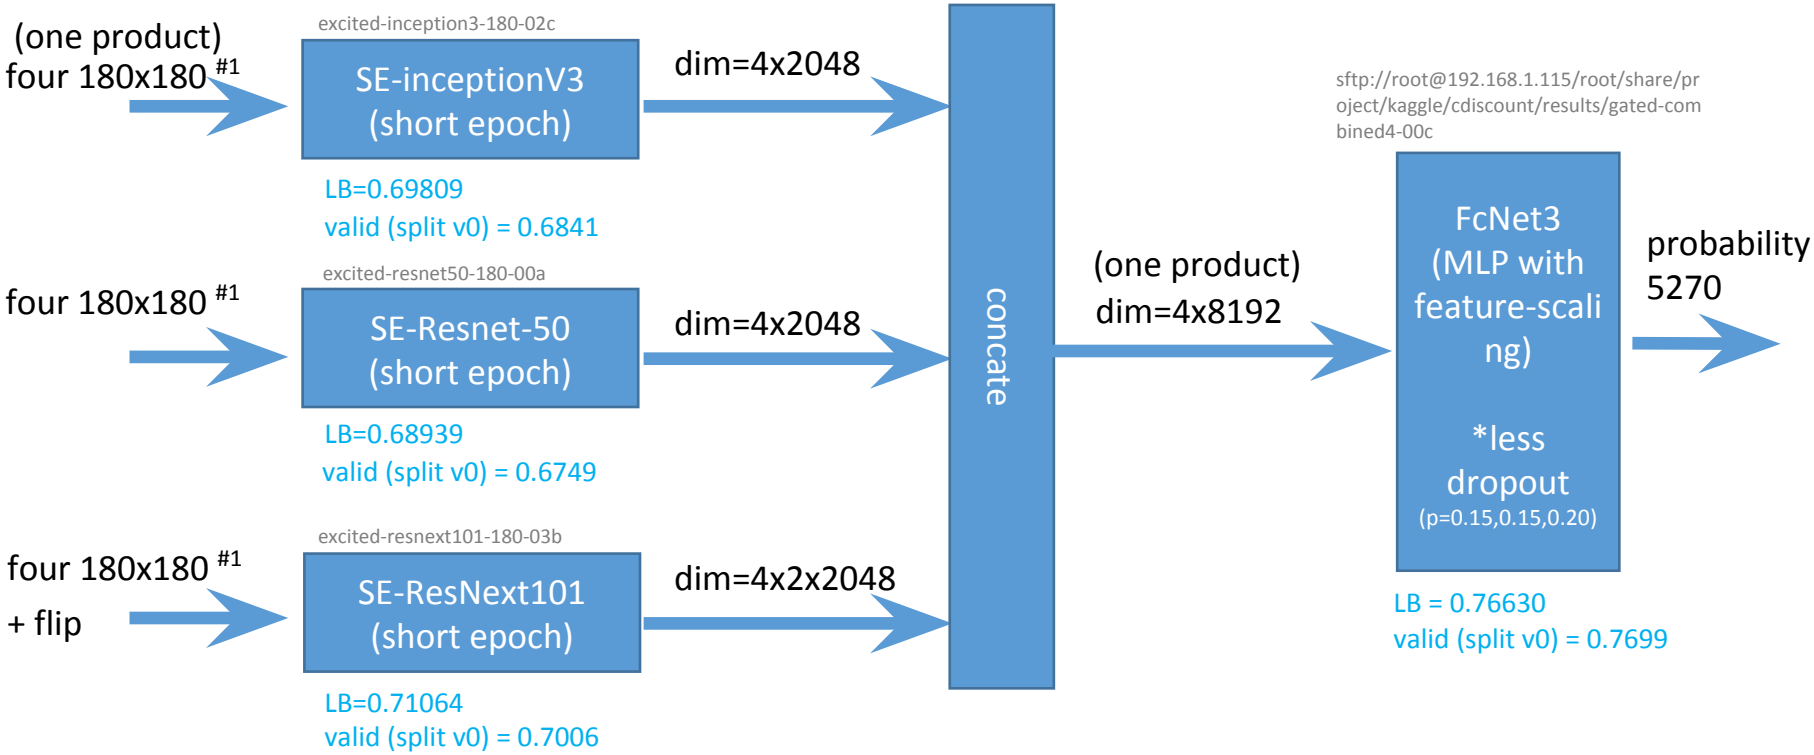

#1 : if the product has less than 4 images, it is stuffed with zero images

[Fuse net]

FcNet3: view feature scaling (gating)

```python
class FcNet3(nn.Module):
    def __init__(self, in_shape=8192, num_classes=5270 ):
        super(FcNet3, self).__init__()
        self.num_classes = num_classes
        in_channels = in_shape

        self.scale   = SEScale(in_channels, in_channels//2)
        self.linear1 = nn.Linear(in_channels, 7168)
        self.relu1   = nn.PReLU()
        self.linear2 = nn.Linear(7168, 4096)
        self.relu2   = nn.PReLU()
        self.fc = nn.Linear(4096, num_classes)


    def forward(self, x):
        #
                              #; print('input ' ,x.size())
        N,V,C = x.size()
        x = F.dropout(x, p=0.50, training=self.training)
        x = self.scale(x) * x
        x = x.sum(dim=1)
        x = self.linear1(x)  #; print('linear1 ',x.size())
        x = self.relu1  (x)
        x = F.dropout(x, p=0.10, training=self.training)
        x = self.linear2(x)  #; print('linear2 ',x.size())
        x = self.relu2  (x)
        x = F.dropout(x, p=0.10, training=self.training)

        x = self.fc(x)
        return x  #logits
```

gate

```python
##-----------------------------------------------------------------
class SEScale(nn.Module):
    def __init__(self, channel, reduction=16):
        super(SEScale, self).__init__()
        self.fc1 = nn.Linear(channel,   reduction)
        self.fc2 = nn.Linear(reduction, channel)

    def forward(self, x):
        x = self.fc1(x)
        x = F.relu(x, inplace=True)
        x = self.fc2(x)
        x = F.sigmoid(x)
        return x
```

N = batch size = 4096
V = num of images per product = 4
C = concat feature dim = 8192

combine.1

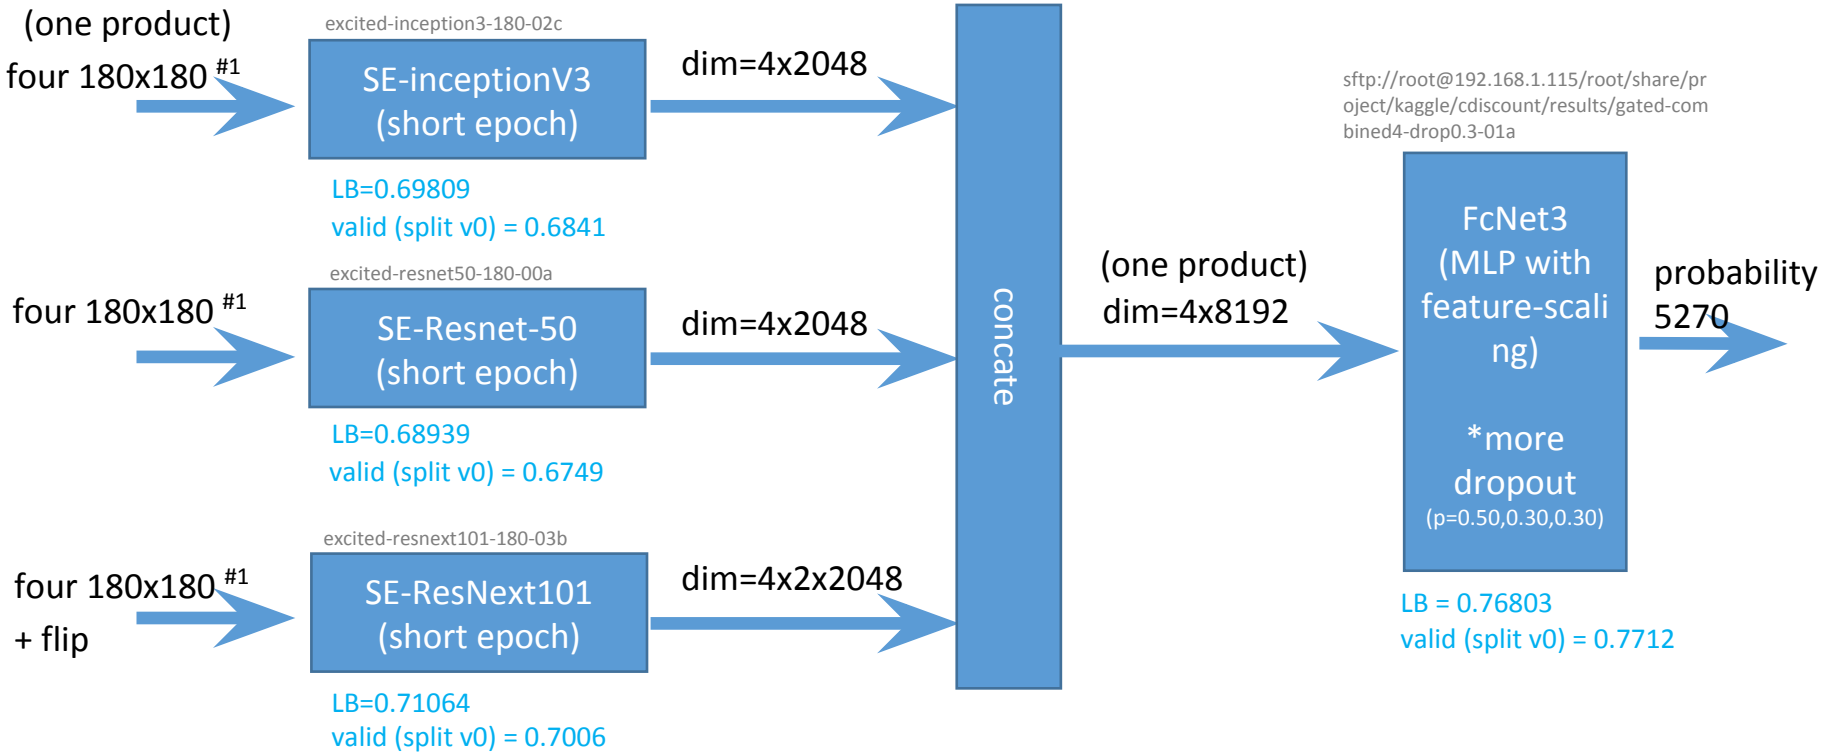

#1 : if the product has less than 4 images, it is stuffed with zero images

combine.2

sftp://root@192.168.1.115/root/share/project/kaggle/cdiscount/results/gated-combined4-05a

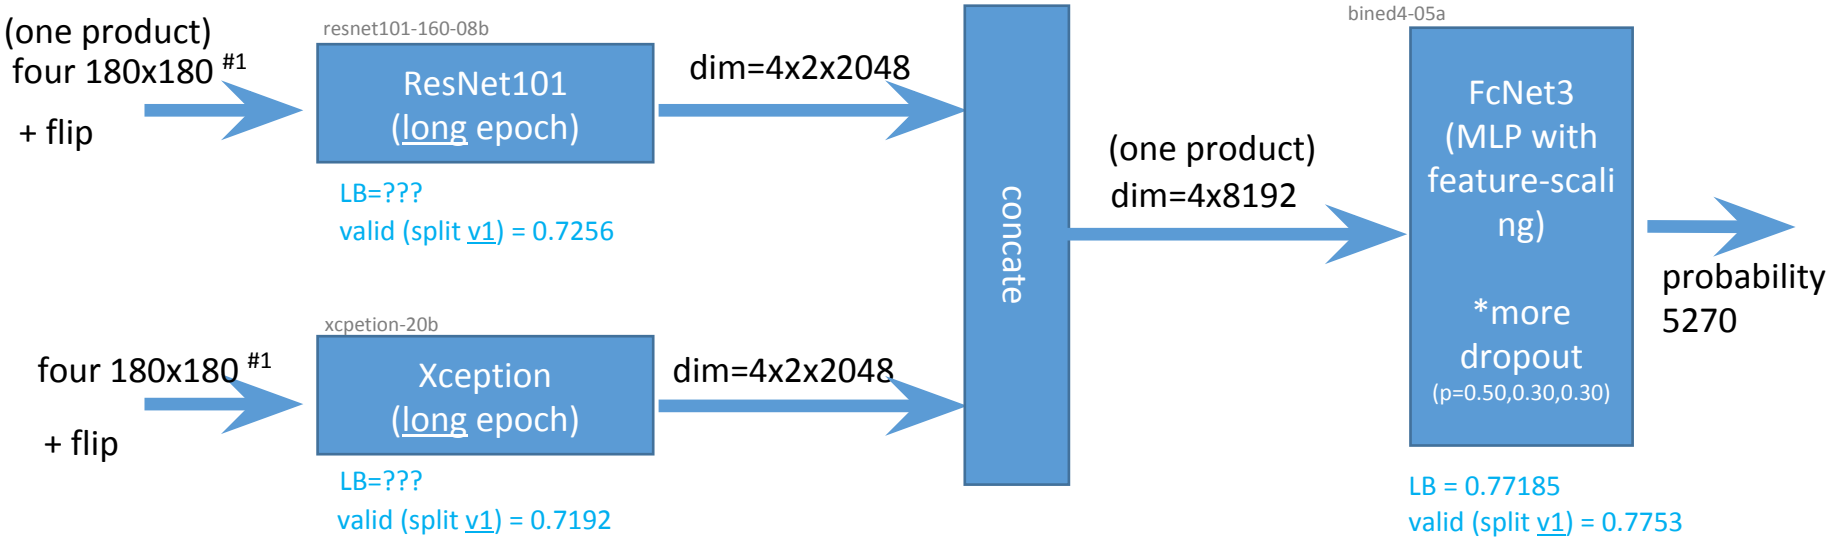

#1 : if the product has less than 4 images, it is stuffed with zero images

combine.5

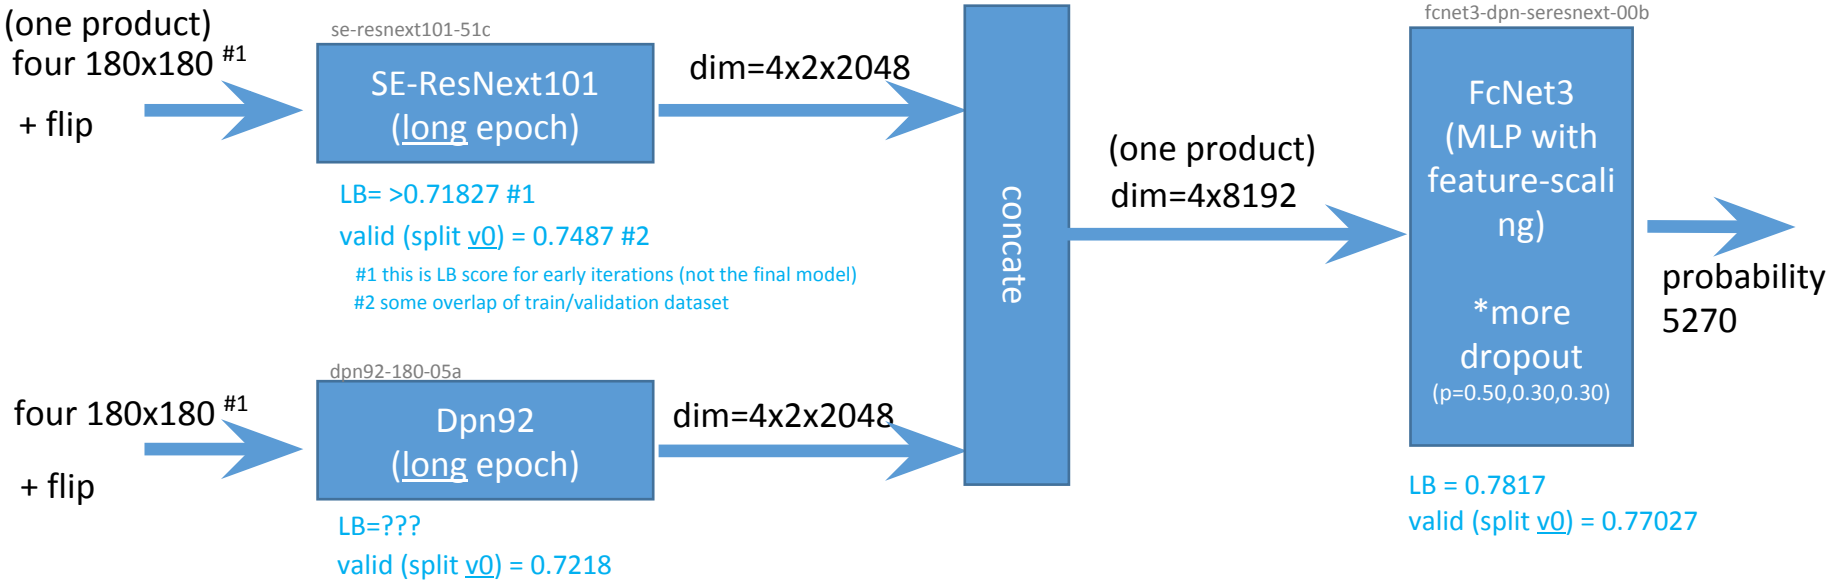

#1 : if the product has less than 4 images, it is stuffed with zero images

combine.5max

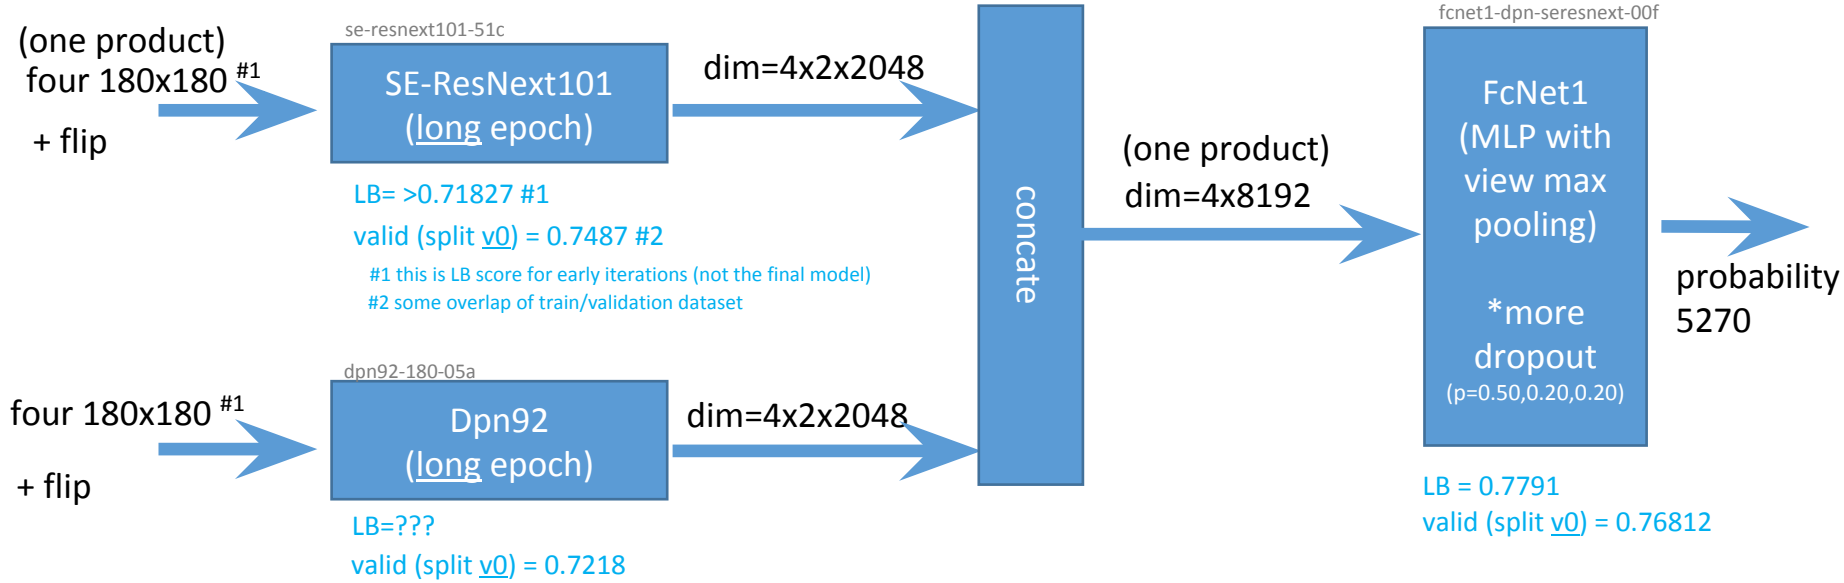

#1 : if the product has less than 4 images, it is stuffed with zero images

combine.5mix

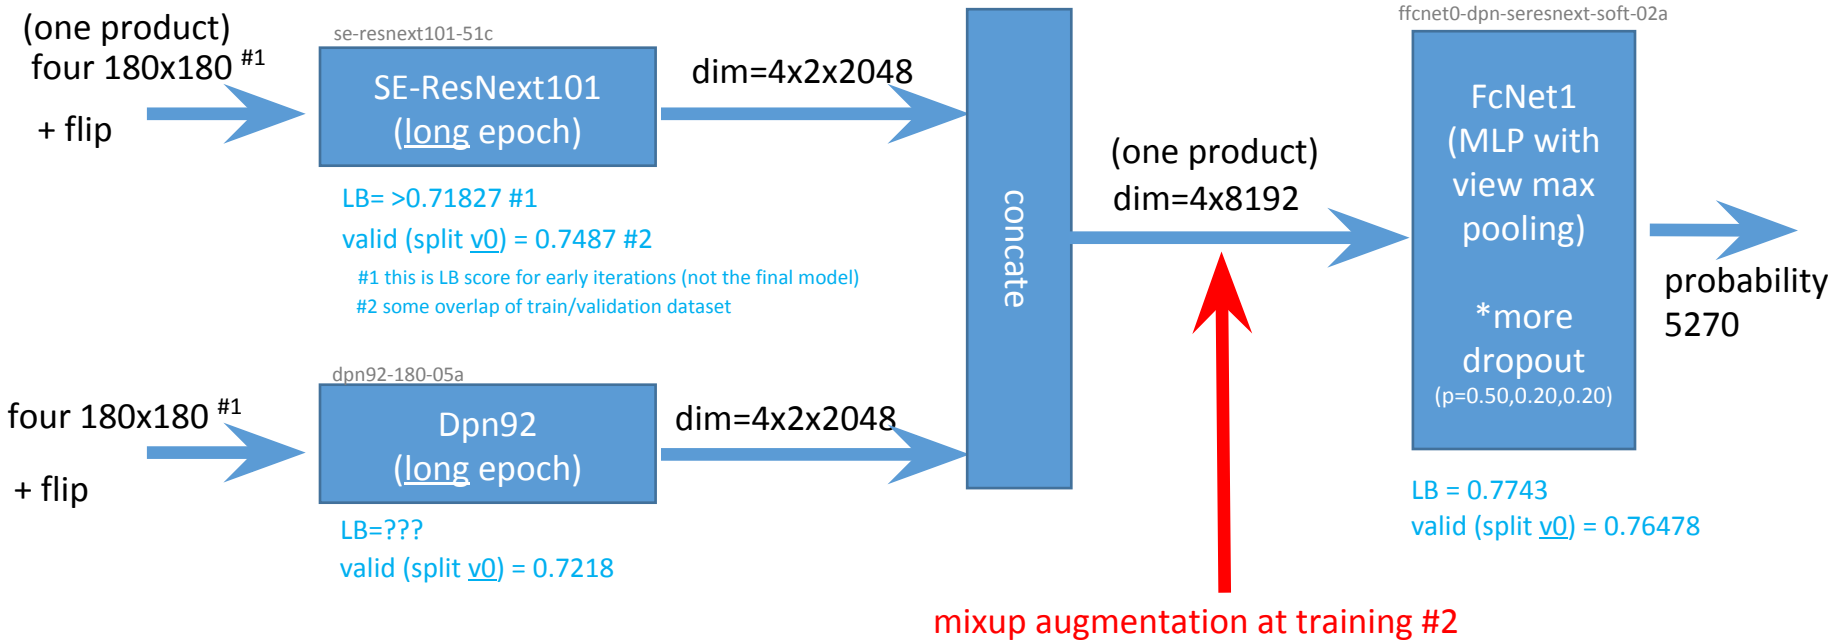

#1 : if the product has less than 4 images, it is stuffed with zero images

#2 : "mixup: Beyond Empirical Risk Minimization" - Hongyi Zhang, ... arvix 2017 (https://arxiv.org/abs/1710.09412)

Azat Davletshin(8th place)

- Finetuned the pre-trained SE-ResNet-50 for about 7 epochs on 161x161 with random crops and horizontal flips
- Found 5 closest (by cosine) pictures from the train for each class
- scores[pic, label] = np.power((1 + search_scores[label, pic, :]) / 2, 35).sum() Got 0.752LB
- Finally: SE-ResNet-101(15 ep) + KNN + SE-ResNet-50 + TTA. Got 0.77883LB

n01z3(9th place)

- Train FC adam (LR 0.001-> 0.0001)

  Polishing, only FC SGD (LR 0.001 -> 0.0001)
  Train full net SGD (LR 0.001-> 0.0001)
  Manual annealing SGD (LR 0.01-> 0.001 -> 0.0001) 1-2 epochs at each stage, repeat twice
  Train without augs SGD LR 0.0001
- Resnext-101, Resnext-50, SE-Resnext-50, Resnet-101, Resnet-152 with 160 input size and TTA10
- Random sample from the item + hard negative sampling

Vladimir Iglovikov(12th place)

41

Single model performance

Vladimir Iglovikov 3 months ago

last comment by steelrose 2mo ago

91

- Resnet 50, 101 and 152, on 160x160 crops, dropping learning rate on the plateau. At the first epoch, all layers except last are frozen
- TTA + geometric mean by model had 0.75 on LB
- Geometric mean of the previous step had 0.77 on LB

Reference

- A summary about our solution:
  https://www.kaggle.com/c/cdiscount-image-classification-challenge/discussion/45733

- The winner’s approach:
  https://www.kaggle.com/c/cdiscount-image-classification-challenge/discussion/45863

Contacts

Pavel Ostyakov - pavelosta@gmail.com

Aleksey Kharlamov - axcel@me.com