

WillPressLeverForFood at DeepHack.RL



CRAIG SWANSON © WWW.PERSPICUITY.COM

Alexander Guschin, Sergey Korolev, Sergey Ovcharenko,
Sergey Sviridov, Max Kharchenko

Moscow 2017


What is DeepHack



- State-of-art problem
- 7 days of programming and lectures from top notch researchers
- Bleeding edge technologies and developments
- And still hackathon and team challenge

Atari Games

100+ games



OpenAI Gym and Reinforcement learning


- Emulator of environments (including Atari 2600 games)

```
import gym

# Make environment
env = gym.make('CartPole-v0')

# Reset environment
env.reset()
for _ in range(1000):


    # Take a random action
    action = env.action_space.sample()
    # Collect reward and new observation
    observation, reward, done, info = env.step(action)
```



Markov Decision Process

Setup


- Set of states: \mathcal{S}
- Set of actions: \mathcal{A}
- Immediate reward: $R_a(s, s')$
- Transitions: $P_a(s, s') = \Pr(s_{t+1} = s' \mid s_t = s, a_t = a)$
- Discount factor: $\gamma \in [0, 1]$



Goal

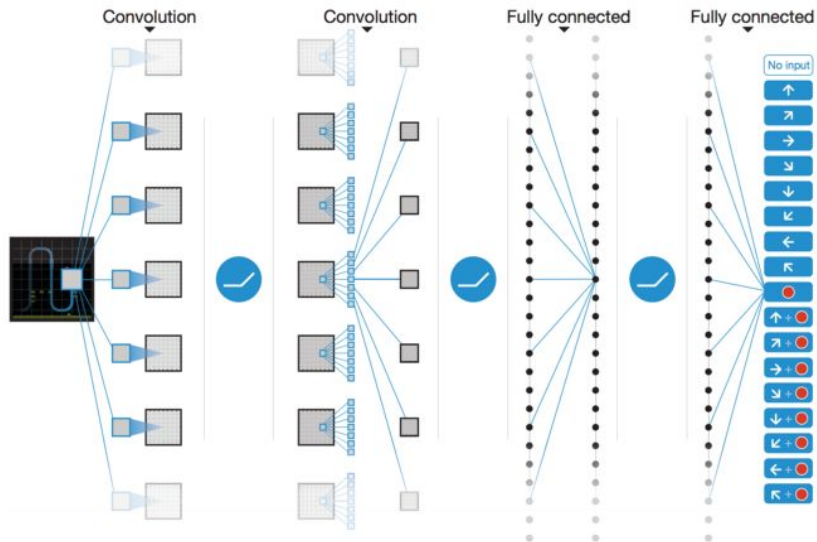
- Find a policy that maximizes the expected future rewards

Q-learning



$$Q(S, A) \leftarrow Q(S, A) + \alpha \left(R + \gamma \max_{a'} Q(S', a') - Q(S, A) \right)$$


Deep Q-Networks




$$Q^*(s, a) \approx \bar{Q}(s, a; \theta)$$

$$L_i(\theta_i) = \mathbb{E} \left(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right)^2$$

Double DQN and Dueling DQN



Natural DQN



Dueling Network:
decomposing state value V and
Action Advantage value A

$$A^{\pi}(s, a) = Q^{\pi}(s, a) - V^{\pi}(s).$$

Policy gradient methods

- ▶ Represent policy by deep network with weights \mathbf{u}


$$a = \pi(a|s, \mathbf{u}) \text{ or } a = \pi(s, \mathbf{u})$$

- ▶ Define objective function as total discounted reward

$$L(\mathbf{u}) = \mathbb{E} [r_1 + \gamma r_2 + \gamma^2 r_3 + \dots \mid \pi(\cdot, \mathbf{u})]$$

- ▶ Optimise objective end-to-end by SGD
- ▶ i.e. Adjust policy parameters \mathbf{u} to achieve more reward

Actor-critic and a3c



- ▶ Actor is updated towards target

$$\frac{\partial l_u}{\partial \mathbf{u}} = \frac{\partial \log \pi(a_t | s_t, \mathbf{u})}{\partial \mathbf{u}} (q_t - V(s_t, \mathbf{v}))$$

- ▶ Critic is updated to minimise MSE w.r.t. target

$$l_v = (q_t - V(s_t, \mathbf{v}))^2$$


- Each worker has own copy of environment
- During rollout worker accumulate gradients
- After rollout worker update global network

Deephack Games




Preliminary stage

Skiing-v0



Main stage


MsPacman-v0




Code Freeze stage

Centipede-v0


Expected behaviour



Reality:Skiing

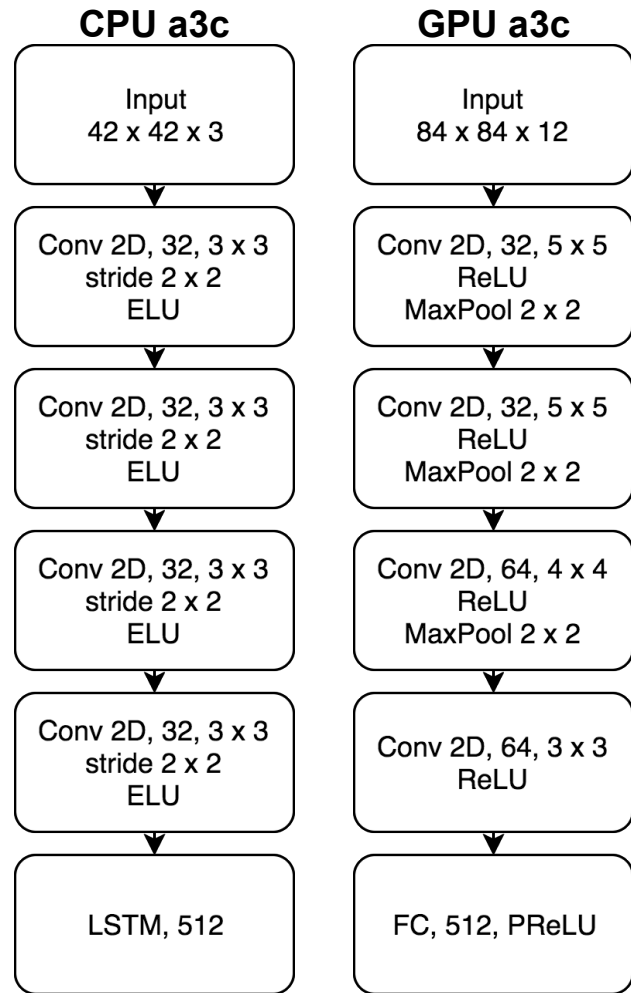


Reality: Pacman



Models selection

- DQN-networks were bad :(
- Two a3c models
 - GPU implementations w/o LSTM
 - CPU implementation with LSTM
- Random actor as backup strategy
- Metaestimator for model selection (by max score)




Hyperparameters optimization

- Preprocessing tweaking
 - A3C preprocessed images to grayscale - made it RGB
 - A3C cropped images and deleted some parts on labyrinth in Pacman
- Gamma tuning $\gamma \in [0, 1]$

Nagibator style: Skiing



Nagibator style: Pacman



Results

ЛУЧШИЙ РЕЙТИНГ ПО 3-М ИГРАМ

	Skiing	Ms. Pac-man	Centepide	Avrg.Rating
WPLFF	-3493	5512	7512	0.790
5vision	-4489	2716	17130	0.753
w0rms	-8493	2311	20328	0.680
bad_skiers_evolved	-8387	2352	17294	0.636
ModestKoalaGrad	-9013	1715	17130	0.579
sw1sh	-16699	2522	5684	0.246
Generation Gap	-9013	-	-	0.194
miptcap	-9013	-	-	0.194
State of the Art	-10853	7534	6297	
	Prioritized DQN	PGQ	Gorila	

Summary

- Try simple methods first
- Don't be afraid to use already existing solutions
- But don't think that these solutions have no mistakes
- Try to dig deeper into model and understand why it performs one way or another

References

- <https://github.com/libfun/deephack3> - our solution
- <https://universe.openai.com/envs> - OpenAI environments
- <https://github.com/openai/universe-starter-agent> - CPU a3c agent with LSTM
- <https://medium.com/emergent-future/simple-reinforcement-learning-with-tensorflow-part-0-q-learning-with-tables-and-neural-networks-d195264329d0#.4rukzup4m> - RL tutorials with TensorFlow
- <https://github.com/ppwwyyxx/tensorpack/tree/master/examples/A3C-Gym> - GPU a3c without LSTM
- <https://arxiv.org/pdf/1602.01783.pdf> - Asynchronous Methods for Deep Reinforcement Learning