

Kaggle TGS Salt Identification Challenge

Sosin Ivan
Kaspersky Lab



TGS SALT IDENTIFICATION CHALLENGE

Segment salt deposits beneath the Earth's surface

Salts and Spices: 11 / 3234


Team members:

- Ivan Sosin



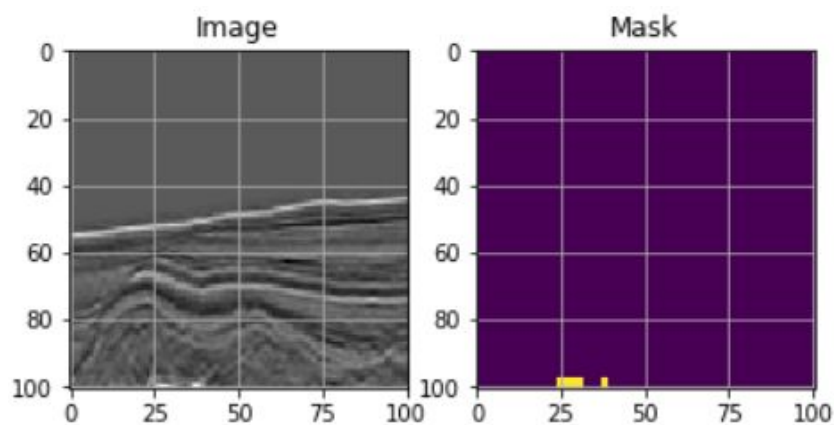

Kaspersky Lab, ML Engineer

- Alexey Rozhkov

Mati, Computer Vision Engineer




Task - binary segmentation



Data orientation

This kind of axes orientation plays a great role in data augmentation



Metric

$$IoU(A, B) = \frac{A \cap B}{A \cup B}.$$

thresholds = (0.5, 0.55, 0.6, 0.65, 0.7, 0.75, 0.8, 0.85, 0.9, 0.95)

$$\frac{1}{|thresholds|} \sum_t \frac{TP(t)}{TP(t) + FP(t) + FN(t)}.$$

Dataset

Train: 4000 images with masks



Test: 18000 images

Train and Test image depths

Image resolution: $101 * 101$

Facts:

- About of 39% of train masks were empty
- There was no full salt masks
- Some masks were just straight lines




Augmentations

- 1) Random horizontal flip
- 2) Random color invert
- 3) Random cutout
- 4) Random gamma-correction
- 5) Random fixed-size crop
- 6) Random crop and scale
- 7) Random rotation up to 10 degrees

Model

Baseline: TernaUSNet
(variation of UNet
architecture)



<https://github.com/ternaus/TernaUSNet>

Tribute to Heng CherKeng

82



Deep semi-supervised learning

[Heng CherKeng](#) 2 months ago

26



how to train/finetune properly for TGS data

[Heng CherKeng](#) 2 months ago

143



common tricks in kaggle image segmentation problems

[Heng CherKeng](#) 2 months ago

27



some baseline performances and observations

[Heng CherKeng](#) 2 months ago

28



using GAN (generative adversarial net) ...

[Heng CherKeng](#) 2 months ago

20



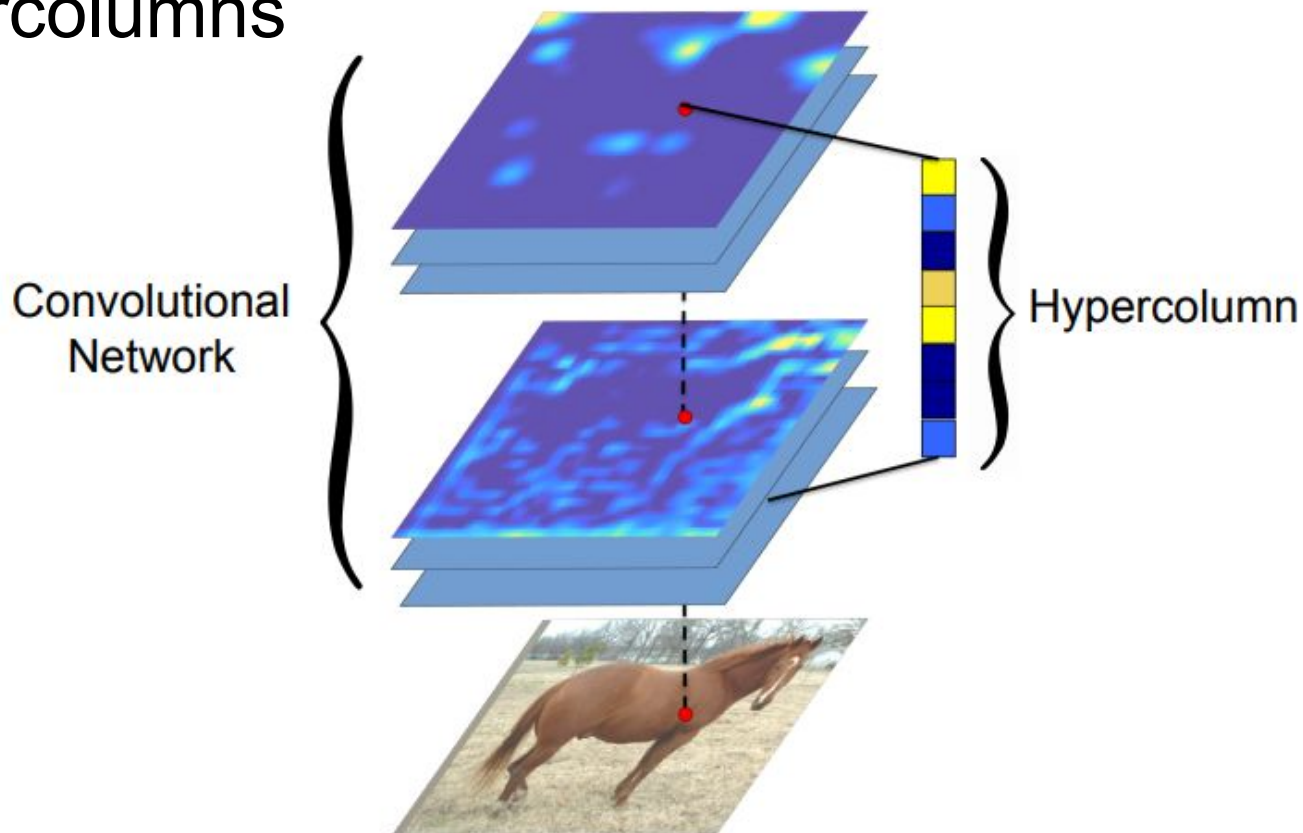
results for deeper models (se-resnext, dpn, densenet, etc)

[Heng CherKeng](#) 2 months ago



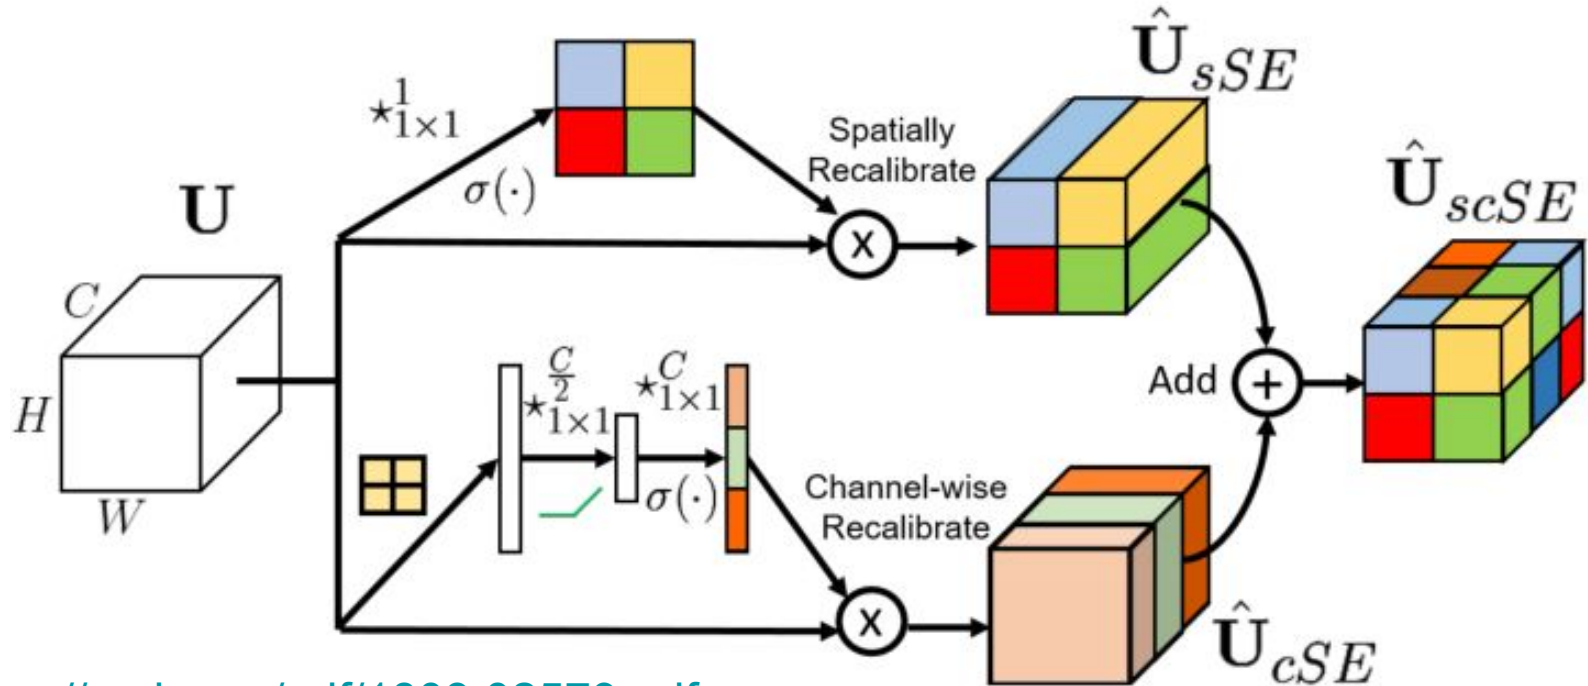
<https://www.kaggle.com/hengck23>

Hypercolumns




<https://arxiv.org/abs/1411.5752>


Concurrent Spatial and Channel Squeeze and Channel Excitation - scSE



Resnext backbone architecture

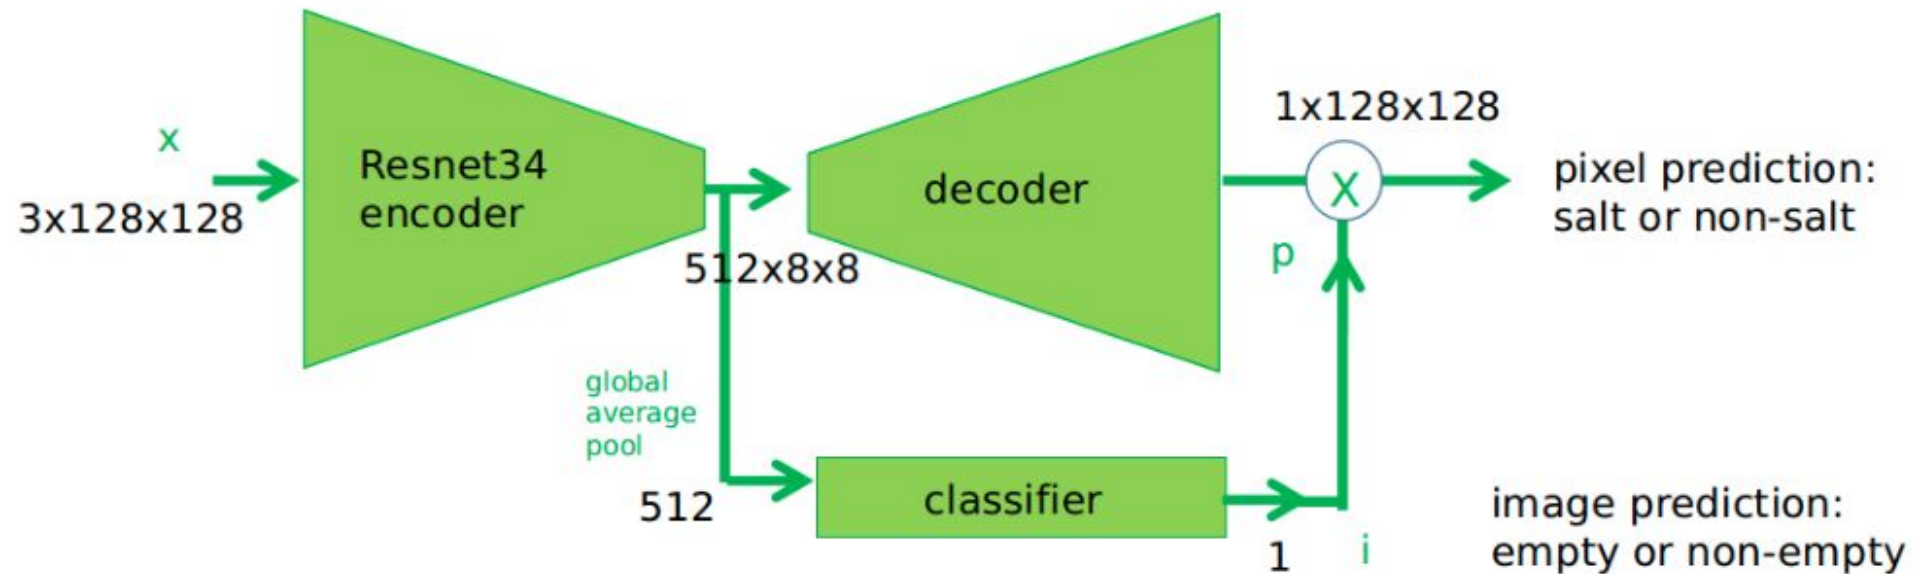


Resnet block

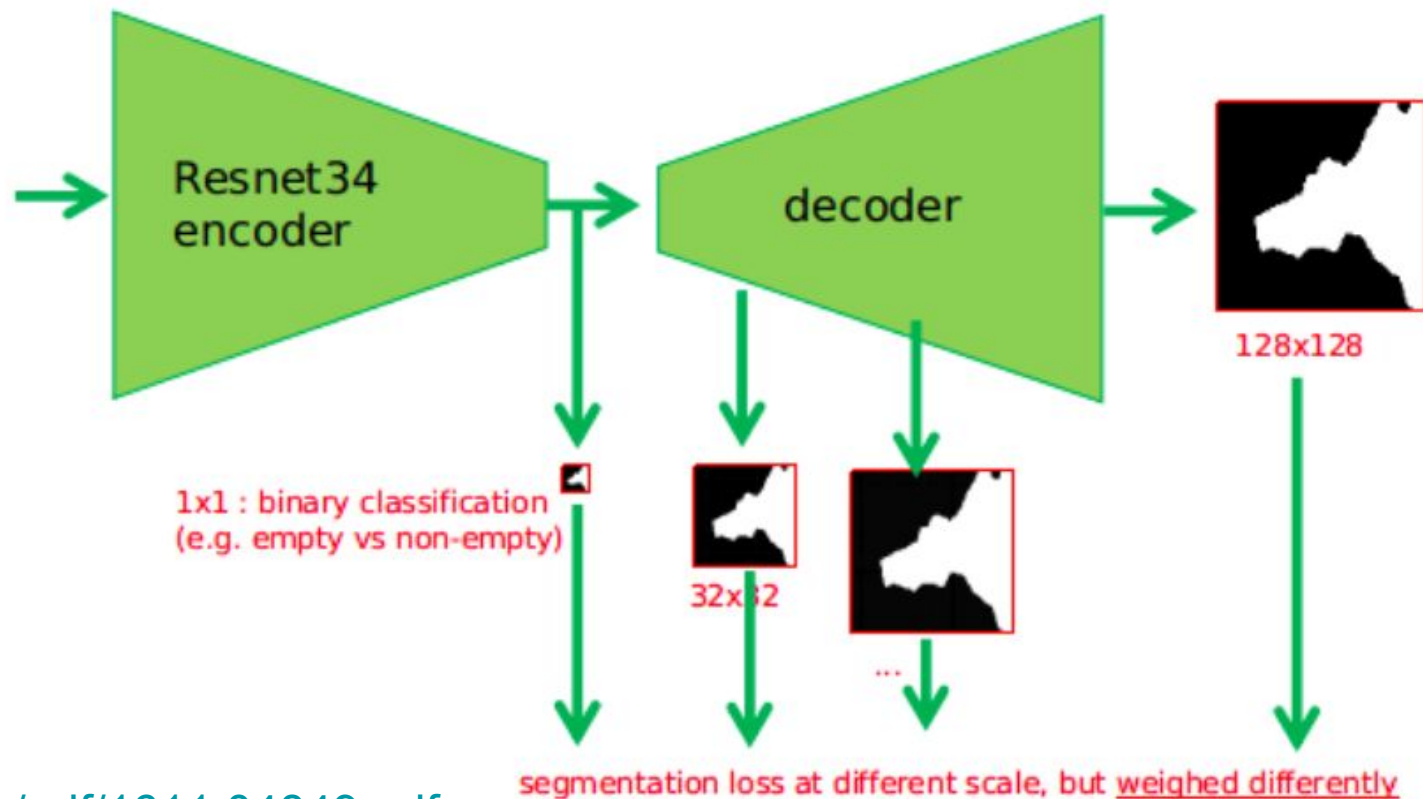


Resnext block with cardinality of 32


Empty / non-empty mask classification



Deeply supervised network




Object Context Network: Self-attention module (SA)



My model architecture

Almost all ideas
blended together

Encoder: pretrained
Se-Resnext-50_32x4d



Decoder features

Ivan:

- Dilated convolutions with dilation from 1 to 5
- Hypercolumns
- OC Module before last Convolution layer
- Deep supervision (zero/nonzero mask, nonzero mask segmentation)

Alexey:

- Spatial and Channel Squeeze Excitation
- Hypercolumns
- Deep supervision (zero/nonzero mask)

Tribute to Peter

194



Improving from 0.78 to 0.84+

Peter 2 months ago

48



+0.01 LB with snapshot ensembling and cyclic lr

Peter 2 months ago

108



Improving from 0.84 to 0.86

Peter a month ago

<https://www.kaggle.com/pestipeti>

Losses

- 1) BCE
- 2) Focal Loss (<https://arxiv.org/pdf/1708.02002.pdf>)
- 3) Dice Loss
- 4) Lovasz loss - a tractable surrogate for the optimization of the intersection-over-union measure in neural networks(<https://arxiv.org/pdf/1705.08790.pdf>)

Training

General:

- SGD: momentum 0.9, weight decay 0.0001
- Batch size: 16
- 5 Fold without stratification

Ivan

- Pretrain for 32 epochs with $\text{lr} = 0.01$ ($0.9 * \text{BCE} + 0.1 * \text{lovasz}$)
- SGDR with cosine annealing ($0.1 * \text{BCE} + 0.9 * \text{lovasz}$):
 - 4 cycles, $\text{lr} = 0.01$, $\text{min_lr} = 0.0001$, $\text{cycle_len} = 64$

Alexey

- SGDR with cosine annealing ($1.0 * \text{BCE} + 0.5 * \text{lovasz}$):
 - 10 cycles, $\text{lr} = 0.05$, $\text{min_lr} = 0.0$, $\text{cycle_len} = 50$

Equipment

Ivan

- Setup: 2x 1080Ti (~11 Gb memory consumption)
- Training time: 8 - 12 hours (per single fold) depending on model modification

Alexey

- Setup: 3x 1080 (~6.6 Gb memory consumption)
- Training time: 16 hours (per single fold)

Ensembling

Ivan:

- LB 0.859 single model, public LB 0.865 5 fold


Alexey:

- LB 0.865 single model, public LB 0.871 5 fold

Ensemble:

- Without post processing: public LB 0.875
- With post processing: public LB 0.881


Mosaics: ~5000 images out of 18000



Post processing

- For almost-vertical ground truth tiles - extrapolate last row downwards till the end of mosaic
- For tiles below almost-vertical predicted tiles - extract downward
- For “caps” in ground truth tiles - zero all tiles below
- For tiles below predicted “caps” - do the same as for ground truth case
- For tiles vertically “sandwiched” between predicted or ground truth masks - replace mask with last row of tile above

Depth feature - 0




Depth feature - 1

“The idea of jigsaw puzzles work great for natural images for the reason that it implicitly exploits continuity of edges within images. Unfortunately, seismic images don't share such properties. Two images from different areas of the world (with some constraints that it is a topic for another post!) can be made to look adjacent/similar under a similar metric. This is because they don't have pronounced discontinuities that make it seem adjacent.”

Questions for organizers

- 1) Why did train and test data come from the same places?
Or did it?
- 2) Why there was no full salt tiles?
- 3) Why tiles were so small?



ODS achievements: half of gold

1. `b.e.s. & phalanx` `@b.e.s.` 🏆 + \$50,000 на двоих
2. `Tim & Sberbank AI Lab` `@alexryzhkov` `@btbpanda` 🏆 + \$25,000 на троих
3. `[ods.ai] topcoders` `@selim_sef` `@Victor` 🏆 + \$15,000 на двоих
8. `DISK` `@igorkrashenyi` `@dennis.sakva`
9. `DAInamite` `@tugstugi`
11. `[ods.ai] Salts and Spices` `@alexisrozhkov` `@sawseen`
13. `Giba&Heng` `@Giba` 🐸
14. `[ods.ai] Argus` `@romul` `@nikolasent`

1st place solution: the first stage

b.e.s.

- Input: 101 -> resize to 192 -> pad to 224
- Encoder: ResNeXt50

Phalanx

- Input: 101 -> resize to 202 -> pad to 256
- Encoder: **Resnet34**, ResNext50

Memes - 0



oyaroshevskiy 12:08 AM

Screenshot from 2018-10-12 00-07-49.png ▼



19



8



6



n01z3 12:04 AM

mem++.jpg ▼





SpaceDoge 🐶 11:13 PM

перед шейкапом

resnet34.png ▾



1st place solution: the second stage

Confident pseudo labeling: confidence as a percentage of confident pixel predictions ($0.2 < \text{probability} < 0.8$)

b.e.s.

- ResNeXt50 was pretrained on confident pseudolabels; and 5 folds were trained on top of them

Phalanx

- Added 1580 pseudolabels to each of 5 folds and trained the model from scratch

1st place solution: the third stage

All pseudo labels from the second stage were taken

Phalanx

- Added 1580 pseudo labels to each of 5 folds and trained the model from scratch
- 2 different models with:
 - Feature Pyramid Attention (<https://arxiv.org/pdf/1805.10180.pdf>)
 - Global Attention Upsample (<https://arxiv.org/pdf/1805.10180.pdf>)
 - Global Convolutional Network (<https://arxiv.org/pdf/1703.02719.pdf>)

Final model is a blend of 2nd stage and 3rd stage

<https://www.kaggle.com/c/tgs-salt-identification-challenge/discussion/69291>

Other things that worked

- 1) SENet154 and other fat encoders, extremely slow
- 2) Symmetric lovasz loss (with minus sign)
- 3) A lot of models (up to 200)
- 4) Simulated shakeup
- 5) SGDR with lr update per each iteration instead of each epoch

Other things that sometimes worked

- 1) Adam
 - 2) Swa
 - 3) Mosaic based folds
 - 4) Depth stratification
 - 5) Upsample instead of Deconvolution in decoder
 - 6) Deeply supervised network with mask edge prediction
 - 7) Best thresholding by KL divergence between train and test distributions
 - 8) CBAM: Convolutional Block Attention Module
- (<https://arxiv.org/pdf/1807.06521.pdf>)

My solution

<https://github.com/iasawseen/Kaggle-TGS-salt-solution>

Thank you for your attention!